


# Managing customized FPGA accelerators with SDSoC!

Linaro Conference Vancouver, CAN - 19SEP2018



---

# Title: Managing customized FPGA accelerators with SDSoC!

---

- Abstract: Using Xilinx SDSoC and HLS as a model, this presentation, will discuss the merits and abilities of customized hardware accelerators involved with performance aspects of embedded systems. Using a designed and tested accelerator on the Avnet Ultra96, performance, power, development time metrics will be compared. Through comparing these various metrics, a new path forward for hardware accelerated software designs will be shown as a path forward for the embedded space. While this is a presentation, the author welcomes discussion!

# Agenda

---

- SDSoC -- What, Why, Where?
- Why is this Important to me?

-----

- Real Example
- Why is this faster?

-----

- Where do I start with SDSoC?
- What is my design flow?

-----

- Platforms, what are they?
- Additional Resources
- Basic HLS

What is  
SDSoC?

- What is HLS?
  - High Level Synthesis
- Language examples?
  - C, C++ and System C to RTL
- What is RTL?
  - Generally refers to Register Transfer Level design abstraction
- Language examples?
  - VHDL (VHSIC Hardware Description Language)
  - Verilog (IEEE 1364, is a hardware description language (HDL))

- Platform
  - Sandbox of resources used by SDSoC to generate solutions
  - Sometimes referred to as the “Solution Space” container
  - Defined by the target hardware
- Solution Space
  - Based on available resources
  - Same solution can be different based available resources
  - Tool will attempt to meet described performance needs

- PS
  - Processing System as a whole
  - Generally contains MOST hardened IP
- PL
  - Programmable Logic
  - Can be referred to as programmable fabric
  - Sometimes referred to as FPGA
- FPGA
  - Field Programmable Gate Array

# What are the Xilinx Tools?



- Vivado
  - The IDE for all Programmable Logic Design
- Supports
  - Partial Reconfiguration
  - Simulation
  - Logic Analysis
  - System Generator (for DSP)
  - HLS
  - More...
- Details of Editions:
  - <https://www.xilinx.com/products/design-tools/vivado/vivado-webpack.html>




# What are the Xilinx Tools?




- Software Development Kit
  - The base tool for all Processing System Designs
- Supports
  - Embedded Development
  - Kernel, Root File System, U-Boot
  - Xilinx provides a Yocto Build Layer
  - PetaLinux binary image creation
- Details of Features:
  - <https://www.xilinx.com/products/design-tools/software-zone/embedded-computing.html#opensource>



# What is SDSoC ?




# What is SDSoC ?



# What is SDSoC ?

- SDSoC is part of the SDx umbrella
- That umbrella includes:
  - SDSoC
    - Software Defined accelerators for a System on Chip
  - SDAccel
    - Software Defined Accel(erators) (targeting more of a PURE Programmable Logic)
  - SDNet
    - Software Defined Specification Environment for Networking



# What is SDSoC?

---



- SDSoC delivers
  - System level profiling
  - Automated software acceleration in programmable logic
  - Automated system connectivity generation
  - Libraries to speed programming

# What is SDSoC?

---



- In other words
  - It is an advanced tool suite
  - Eases the burden of managing the resources
  - Eases the burden of managing the interconnect of more complex designs
  - Deals with the complexities of tool suites for you!
  - All the while providing a better embedded design experience.

# What is SDSoC?




- So doesn't SDSoC pick up where Vivado HLS leaves off??
  - Isn't this just HLS? No!
  - SDSoC is more a resource manager providing an easier path to using the suite of included tools (calling SDK, HLS, Vivado, etc.)
- What does all this mean?
  - Your job with SDSoC is to create a platform
  - You define what resources are available to the platform
  - You are to identify functions for acceleration (profiling or manually)
  - SDSoC will look at the C/C++/OpenCL and USE the available resources to most optimally hit your power/performance targets!

# One way to think about this...

---

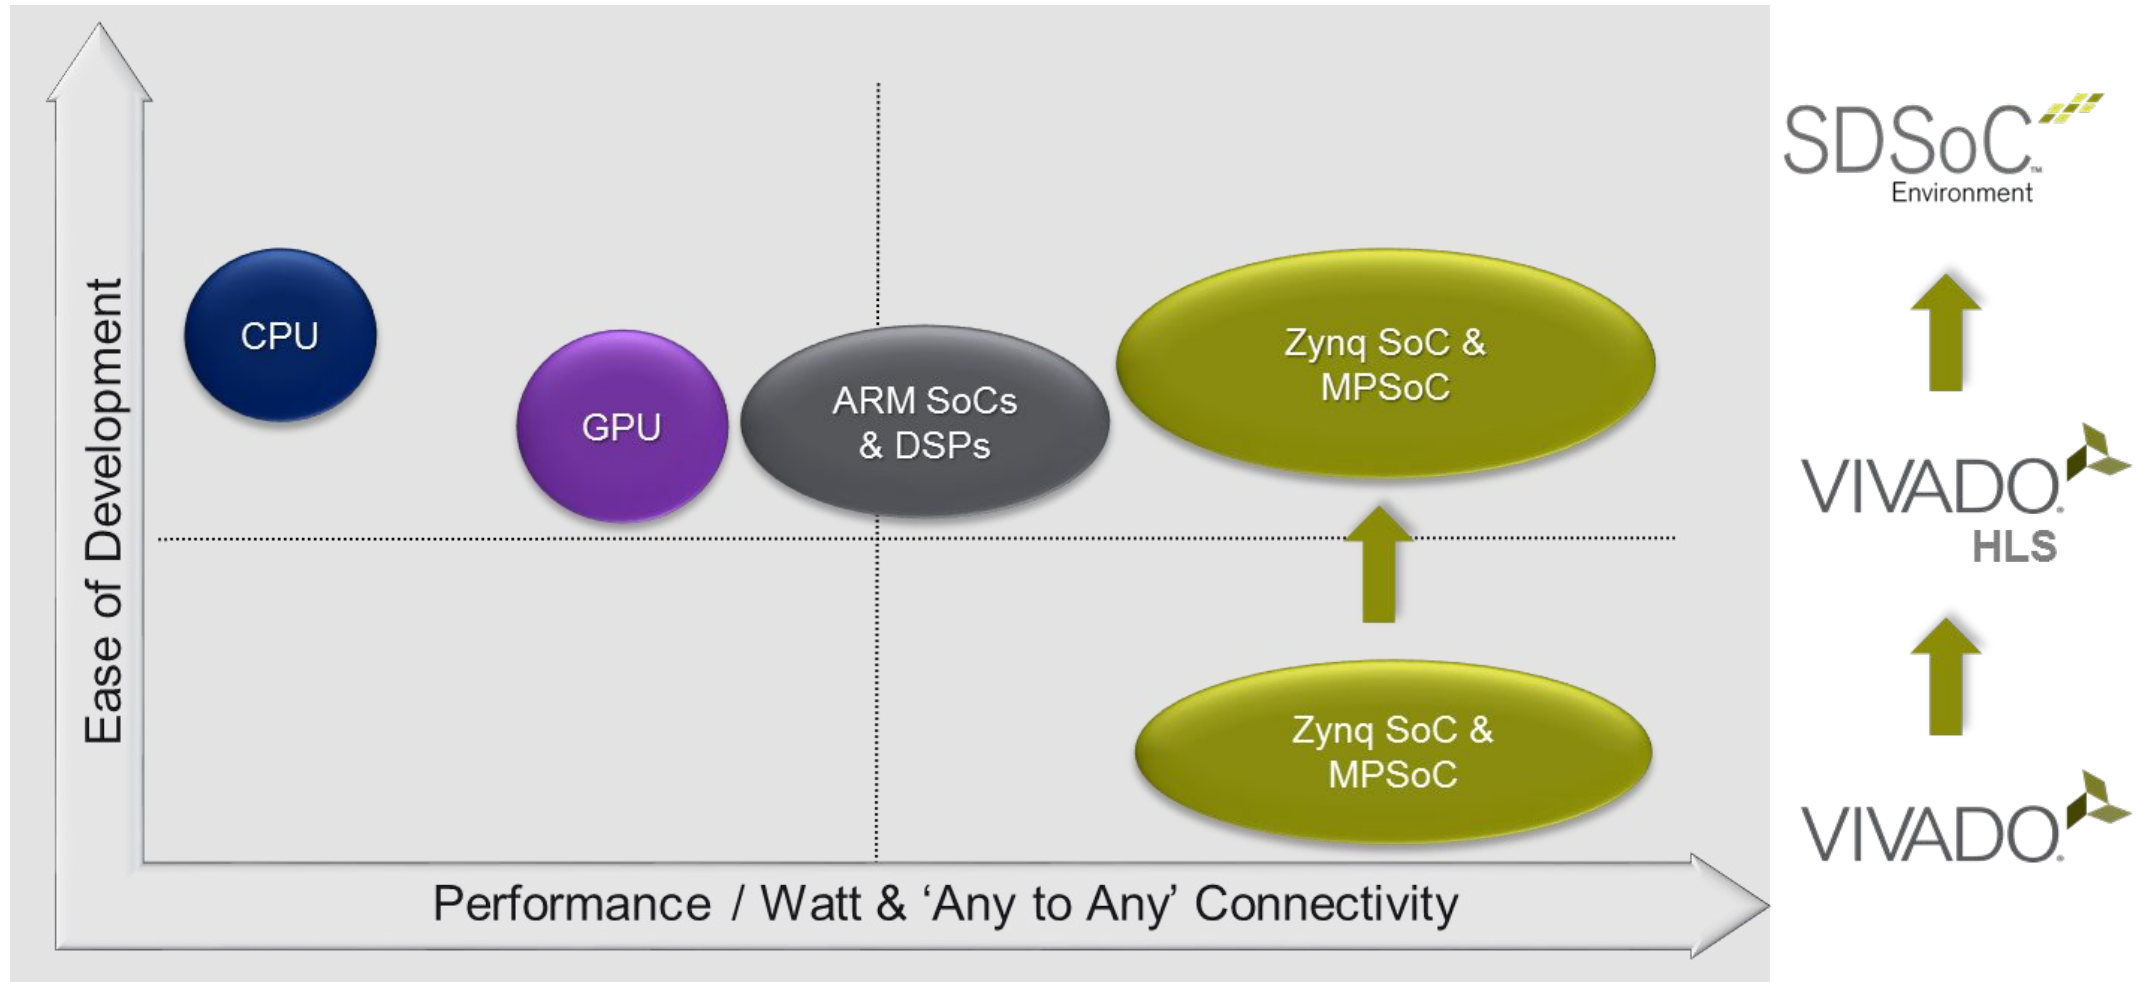
- SDSoC works in layers
- This can be thought of as a pyramid





Why is this  
important to  
me?

# Scaling the Productivity with Technology Advancement



# SDSoC Makes Everyone More Productive

## System Team

- Explore HW-SW partitions and architecture rapidly in C/C++

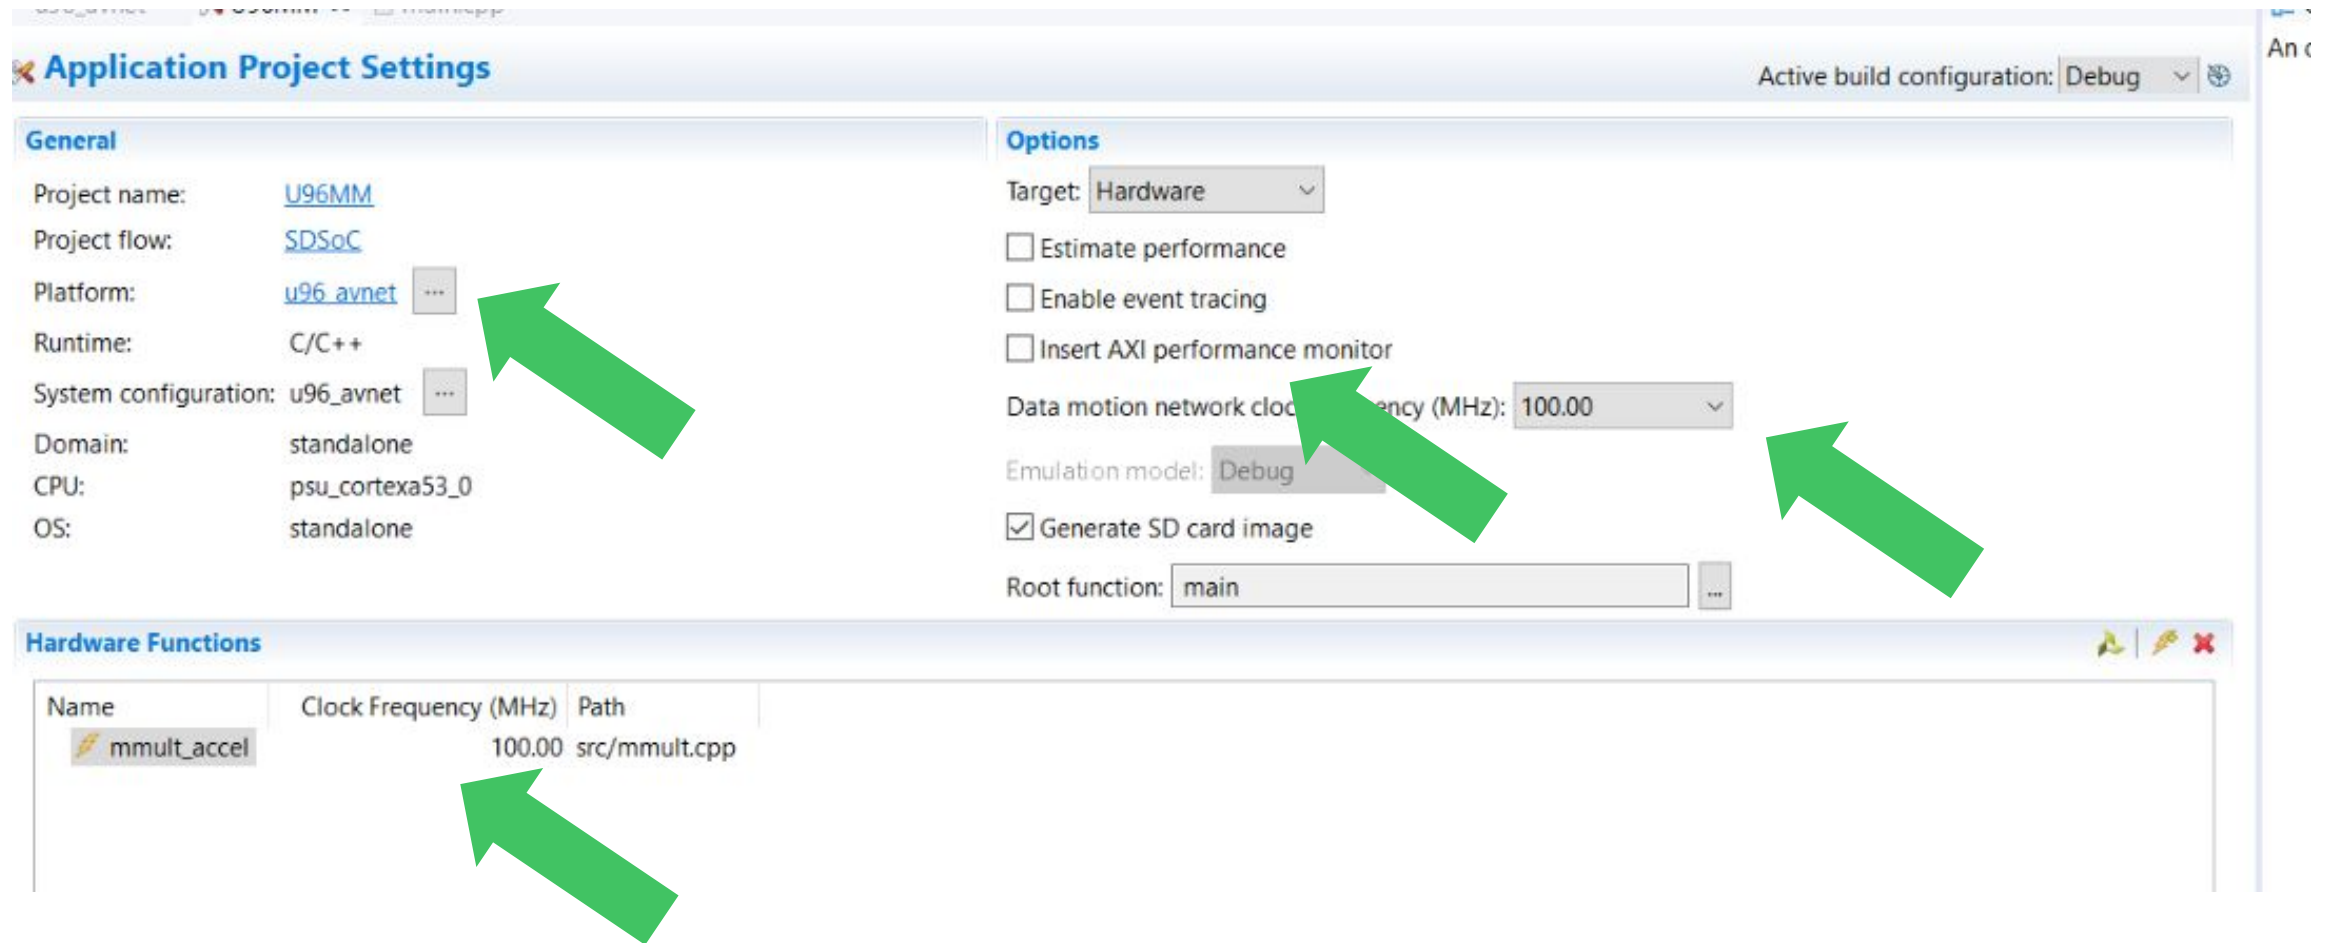
## Hardware Team

- No need to translate algorithms to HDL
- Directly optimize IP on C/C++
- Build IO system into a re-usable platform

## Software Team

- A complete SW development environment for the full Zynq system

# SDSoC Makes Everyone More Productive

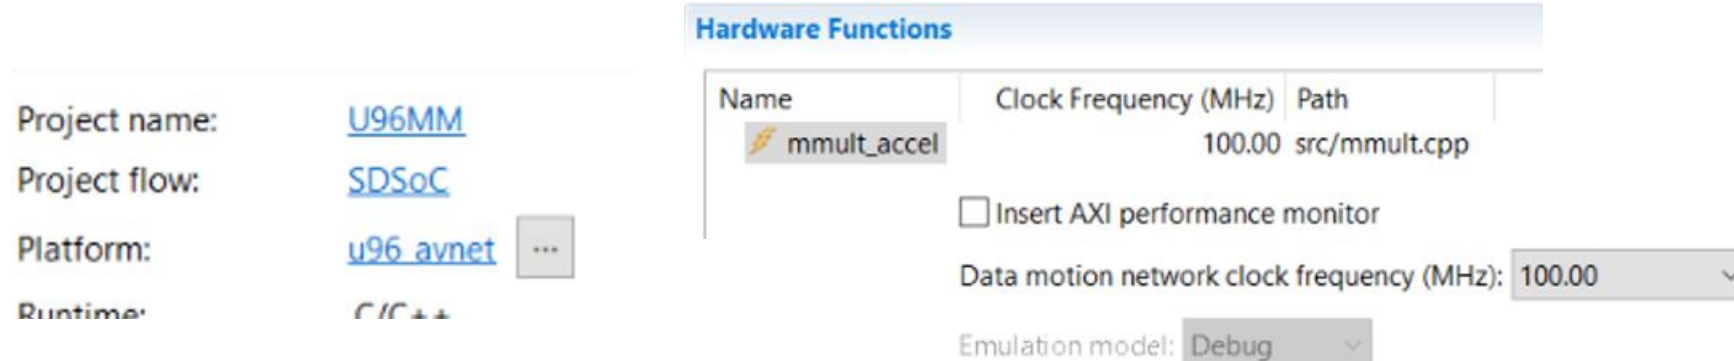


# SDSoC Makes Everyone More Productive

Project name: [U96MM](#)  
Project flow: [SDSoC](#)  
Platform: [u96\\_avnet](#)   
Runtime: [C/C++](#)

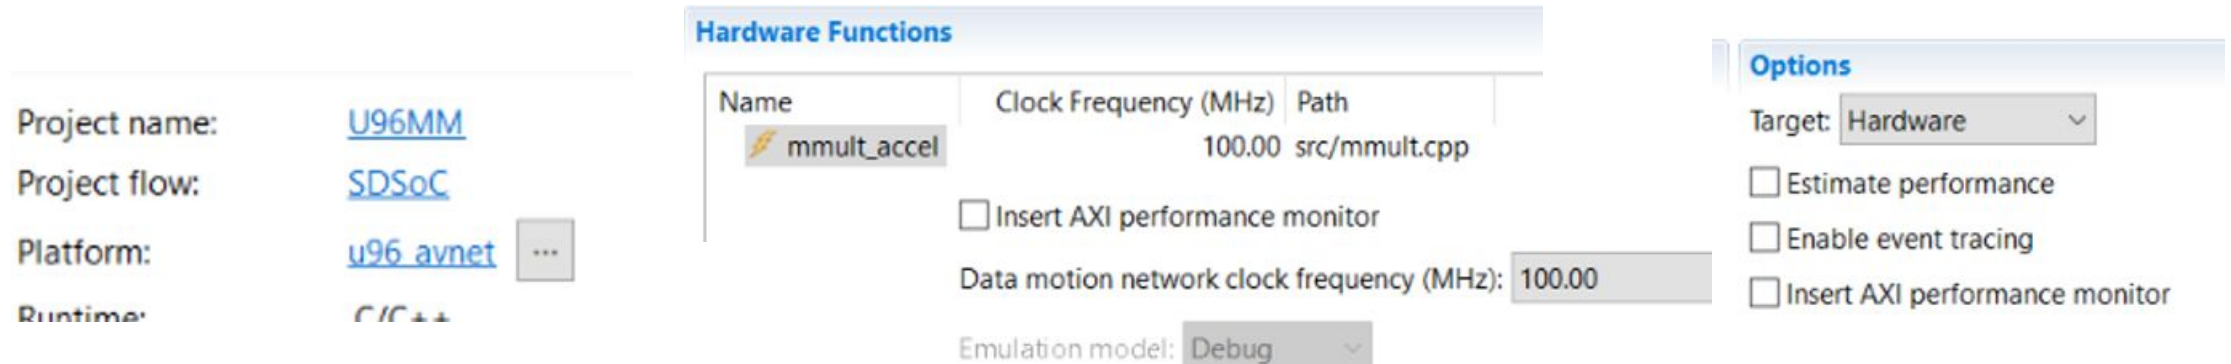
- By working at a system level in the solution space
  - Need only create the solution once
  - You can retarget to a new platform with different resources with a simple retarget

# SDSoC Makes Everyone More Productive



- By working at a system level in the solution space
  - Need only create the solution once
  - You can retarget to a new platform with different resources with a simple retarget
  - You can also quickly and easily change performance behaviors through a few selections

# SDSoC Makes Everyone More Productive



- By working at a system level in the solution space
  - Need only create the solution once
  - You can retarget to a new platform with different resources with a simple retarget
  - You can also quickly and easily change performance behaviors through a few selections
  - Lastly, there is a Performance Estimation tool which will significantly help you target WHICH algorithms should be targeted to hardware acceleration

---

## ...Why is this important to me?

---

- You do not have to be a FPGA expert!
  - SDSoc is another step closer to more easily empowering Engineers who have little to no FPGA/HDL experience
- Problems are more complex
  - SDSoc allows a designer to take a system's approach
  - SDSoc allows you to import your existing Zynq design and start developing new applications in C/C++/OpenCL
  - Designers can concentrate more effort on the solution
    - Instead of determining gate level logic



# Real Example

# How about a REAL example?

```
class perf_counter
{
public:
    uint64_t tot, cnt, calls;
    perf_counter() : tot(0), cnt(0), calls(0) {};
    inline void reset() { tot = cnt = calls = 0; }
    inline void start() { cnt = sds_clock_counter(); calls++; };
    inline void stop() { tot += (sds_clock_counter() - cnt); };
    inline uint64_t avg_cpu_cycles() { return (tot / calls); };
};

sw_ctr.start();
mmult_golden(A, B, C_sw);
sw_ctr.stop();

hw_ctr.start();
mmult_accel(A, B, C);
hw_ctr.stop();
```

# How about a REAL example?

```
void mmult_golden(float *A, float *B, float *C)
{
    for (int row = 0; row < N; row++) {
        for (int col = 0; col < N; col++) {
            float result = 0.0;
            for (int k = 0; k < N; k++) {
                result += A[row*N+k] * B[k*N+col];
            }
            C[row*N+col] = result;
        }
    }
}
```

# How about a REAL example?

```
void mmult_accel(float A[N*N], float B[N*N], float C[N*N])
{
    float _A[N][N], _B[N][N];
    #pragma HLS array_partition variable=_A block factor=8 dim=2
    #pragma HLS array_partition variable=_B block factor=8 dim=1

    for(int i=0; i<N; i++) {
        for(int j=0; j<N; j++) {
            #pragma HLS PIPELINE
            _A[i][j] = A[i * N + j];
            _B[i][j] = B[i * N + j];
        }
    }

    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            #pragma HLS PIPELINE
            float result = 0;
            for (int k = 0; k < N; k++) {
                float term = _A[i][k] * _B[k][j];
                result += term;
            }
            C[i * N + j] = result;
        }
    }
}
```

# How about a REAL example?

```
void mmult_accel(float A[N*N], float B[N*N], float C[N*N])
{
    float _A[N][N], _B[N][N];
    #pragma HLS array_partition variable=_A block factor=8 dim=2
    #pragma HLS array_partition variable=_B block factor=8 dim=1

    for(int i=0; i<N; i++) {
        for(int j=0; j<N; j++) {
            #pragma HLS PIPELINE
            _A[i][j] = A[i * N + j];
            _B[i][j] = B[i * N + j];
        }
    }

    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            #pragma HLS PIPELINE
            float result = 0;
            for (int k = 0; k < N; k++) {
                float term = _A[i][k] * _B[k][j];
                result += term;
            }
            C[i * N + j] = result;
        }
    }
}
```

```
void mmult_golden(float *A, float *B, float *C)
{
    for (int row = 0; row < N; row++) {
        for (int col = 0; col < N; col++) {
            float result = 0.0;
            for (int k = 0; k < N; k++) {
                result += A[row*N+k] * B[k*N+col];
            }
            C[row*N+col] = result;
        }
    }
}
```

# How about some REAL numbers?

---

- When working in the Processor, typically we use **FLOPS**.
  - The A-53 processors in this design are configured for **1200MHz**
- When working in the Programmable Logic, typically we use **CLOCKS**.
  - The Programmable Logic is running at **100MHz, 300MHz, 375MHz**


Why the difference?

PS has many pipelines, caches, word execution times, etc., which can confuse the number of clock cycles

PL runs on an exact cadence, all pipelines, etc. can be exactly modeled as **YOU CREATED THEM!**

# How about some REAL numbers?

- PC: I7- 4 Core 3.17GHz; 48GB Ram; SSD
- Platform Build Times?
  - Vivado 11:30
  - DSA Creation 0:10
  - SDx 4:00
- Solution Build Times?
  - Program Setup 0:33
  - For 100MHz PL 35:34 (1<sup>st</sup> Run)
  - For 100MHz PL 24:34 (2nd Run)
  - For 300MHz PL 35:54
  - For 375MHz PL 42:22





# How about some REAL numbers?

```
COM11:115200baud - Tera Term VT
File Edit Setup Control Window Help
Xilinx Zynq MP First Stage Boot Loader
Release 2018.2 Sep 11 2018 - 23:07:15
Testing 1024 iterations of 32x32 floating point mmult...
Average number of CPU cycles running mmult in software: 1578617
Average number of CPU cycles running mmult in hardware: 65476
Speed up: 24.1099
TEST PASSED

COM11:115200baud - Tera Term VT
File Edit Setup Control Window Help
Xilinx Zynq MP First Stage Boot Loader
Release 2018.2 Sep 11 2018 - 23:07:15
Testing 1024 iterations of 32x32 floating point mmult...
Average number of CPU cycles running mmult in software: 1578611
Average number of CPU cycles running mmult in hardware: 40918
Speed up: 38.5799
TEST PASSED
```

Hardware Functions		
Name	Clock Frequency (MHz)	Path
mmult_accel	100.00	src/mmult.cpp

Hardware Functions		
Name	Clock Frequency (MHz)	Path
mmult_accel	300.00	src/mmult.cpp

Hardware Functions		
Name	Clock Frequency (MHz)	Path
mmult_accel	375.00	src/mmult.cpp

Did not meet timing



# Data – in detail

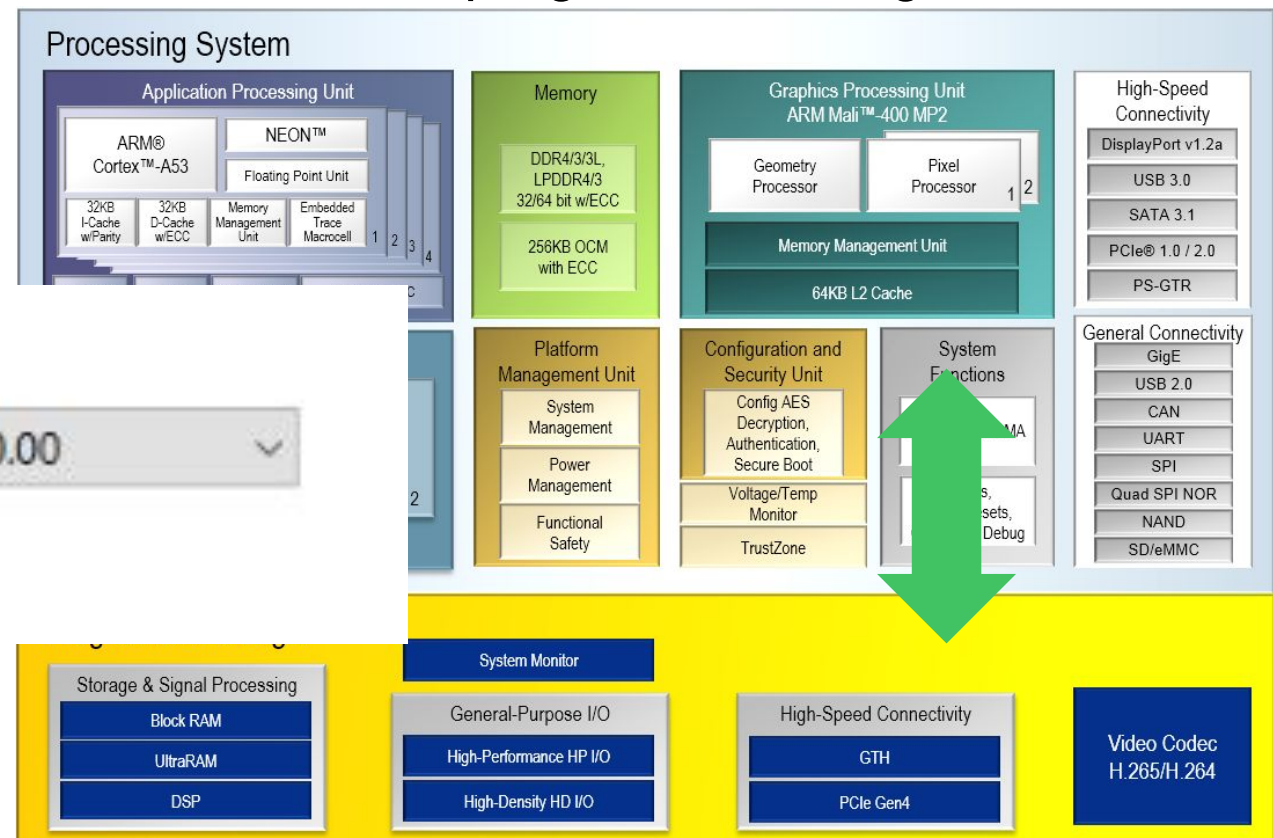
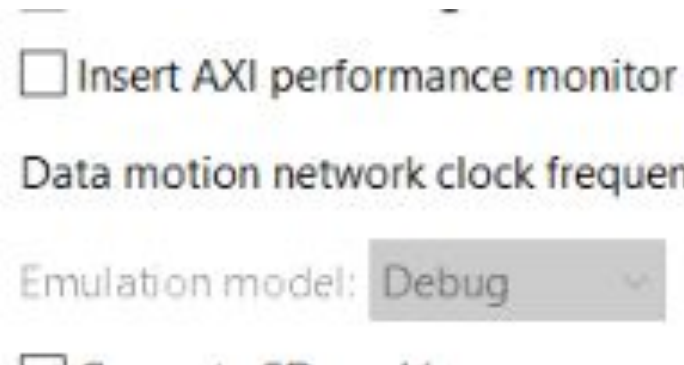
			Acceleration over:			
			PS-Single	PS-Quad	PL 100	PL 300
	MHz	No Clocks	le	ad		
PS Single Core	1200	1578611	1	0.3	0	0
PS Quad Core	4800	394653	4	1	0.2	0.1
PL	100	65476	24.1	6	1	0.6
PL	300	40918	38.6	9.6	1.6	1

Note: RED numbers are theoretical

- To perform 1024 cycles of a 32x32 matrix multiply
  - The PS takes 1,578,611 clocks @ 1200MHz
  - The PL takes 65,476 PS clocks @ 100MHz
  - The PL takes 40,918 PS clocks @ 300MHz
  - 24x increase with PL @ 100MHz
  - ~39x increase with PL @ 300MHz
  - 1.6x increase with PL 100→300MHz

# Why is it not Linear?

- If we process 3x faster, would we not see 3x improvement in our calculation?
- Remember that we have to move the data INTO the programmable logic as well as move that data OUT
- This clock was NOT adjusted during testing



Why is this  
faster?

# Parallel Execution

```
void mmult_accel(float A[N*N], float B[N*N], float C[N*N])
{
    float _A[N][N], _B[N][N];
    #pragma HLS array_partition variable=_A block factor=8 dim=2
    #pragma HLS array_partition variable=_B block factor=8 dim=1

    for(int i=0; i<N; i++) {
        for(int j=0; j<N; j++) {
            #pragma HLS PIPELINE
            _A[i][j] = A[i * N + j];
            _B[i][j] = B[i * N + j];
        }
    }

    for (int i = 0; i < N; i++) {
        for (int j = 0; j < N; j++) {
            #pragma HLS PIPELINE
            float result = 0;
            for (int k = 0; k < N; k++) {
                float term = _A[i][k] * _B[k][j];
                result += term;
            }
            C[i * N + j] = result;
        }
    }
}
```

← Loading Arrays into Programmable Logic

← Programmable Logic Executes in parallel

← Datamover used to place into DDR4

---

# Parallel Execution – Limitations?

---

- By breaking up algorithms into smaller pieces
  - How parallel does your algorithm support?
- Depending on physical limitations
  - power density of processing unit
  - transistor count
  - frequency limitations (just like our 375MHz case)

# Amdahl's Law

- My take
  - Due to losses in efficiency of parallelism as well as limitations on the slowest part, an ever increasing number of processing units is needed to maintain your speedup


Amdahl's law can be formulated in the following way:

$$S_{\text{latency}}(s) = \frac{1}{(1 - p) + \frac{p}{s}}$$

where

- $S_{\text{latency}}$  is the theoretical speedup of the execution of the whole task;
- $s$  is the speedup of the part of the task that benefits from improved system resources;
- $p$  is the proportion of execution time that the part benefiting from improved resources originally occupied.

[https://en.wikipedia.org/wiki/Amdahl%27s\\_law](https://en.wikipedia.org/wiki/Amdahl%27s_law)



# Parallel Execution – Programmable Fabric Shines!

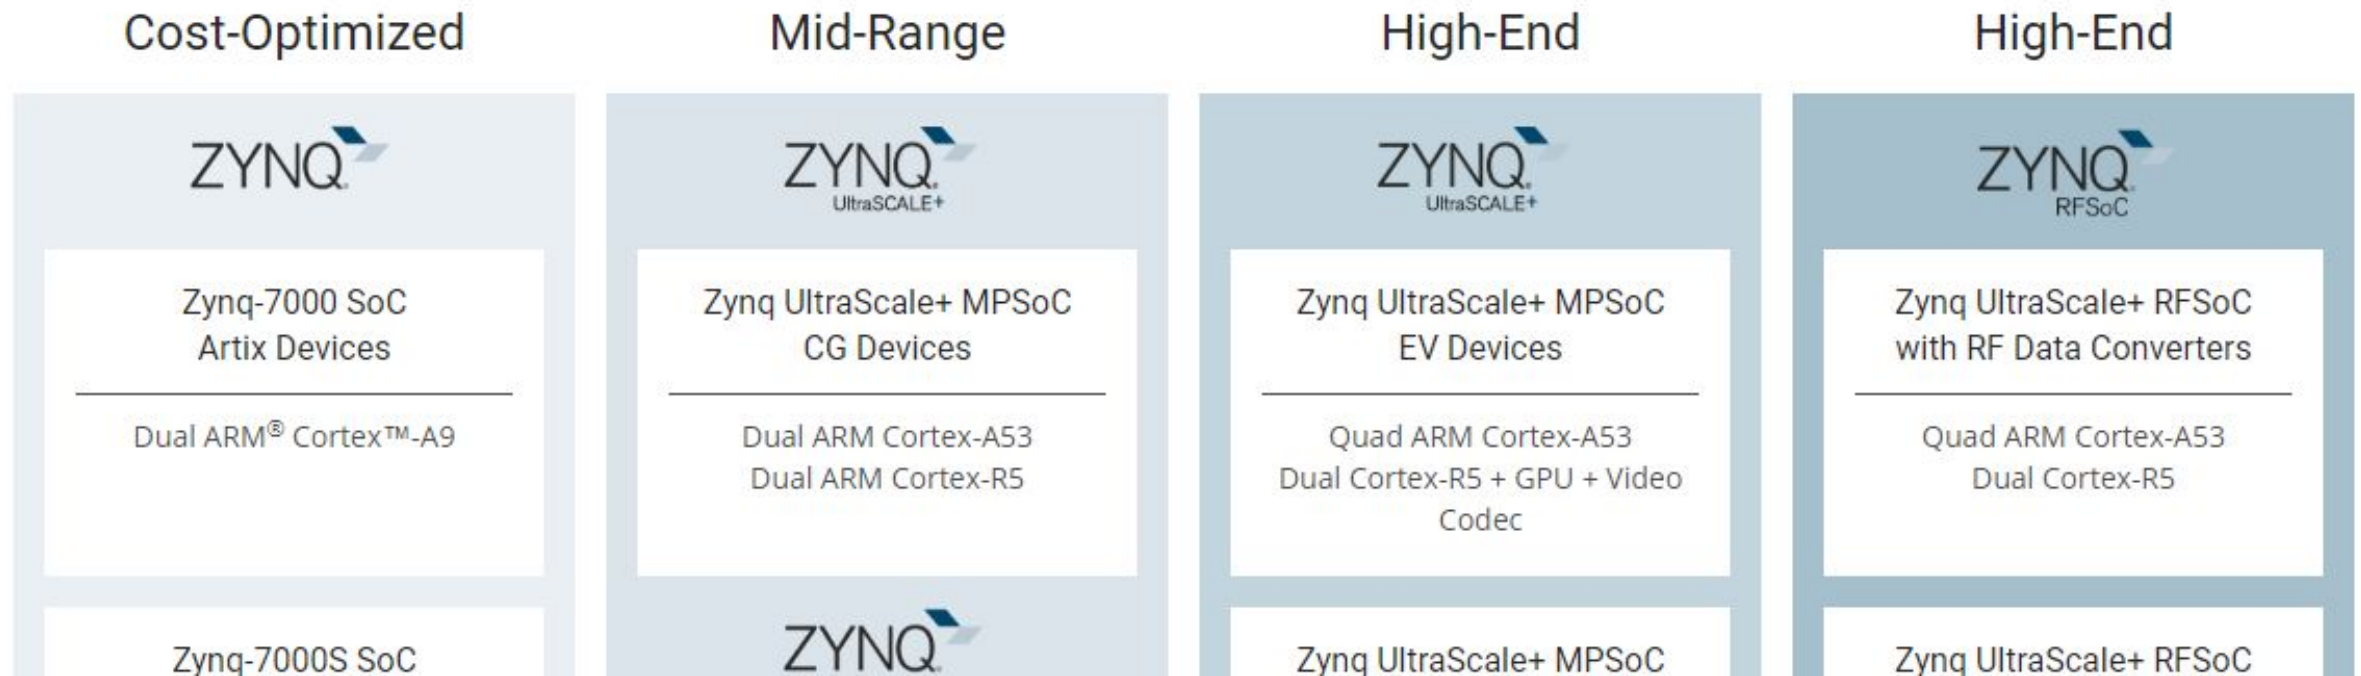
---

- In the case of Programmable Logic
  - We break the problem up into the smallest pieces
  - Create a customized, dedicated accelerator
  - Providing resources specifically to handle the problem
- In the Matrix Multiply
  - Ideally we would have 1 multiplier for EACH multiplication that is needed as well as an adder for EACH resultant, providing a 2 to 3 PL clock result (~20-30 PS clocks)
- In our SDSoC solution
  - The tool evaluated the resources we had at hand and provided enough resources to get us down to ~40 – 64 PS clocks for each matrix multiply!
- Compare this to the PS
  - ~1542 PS clocks per matrix multiply

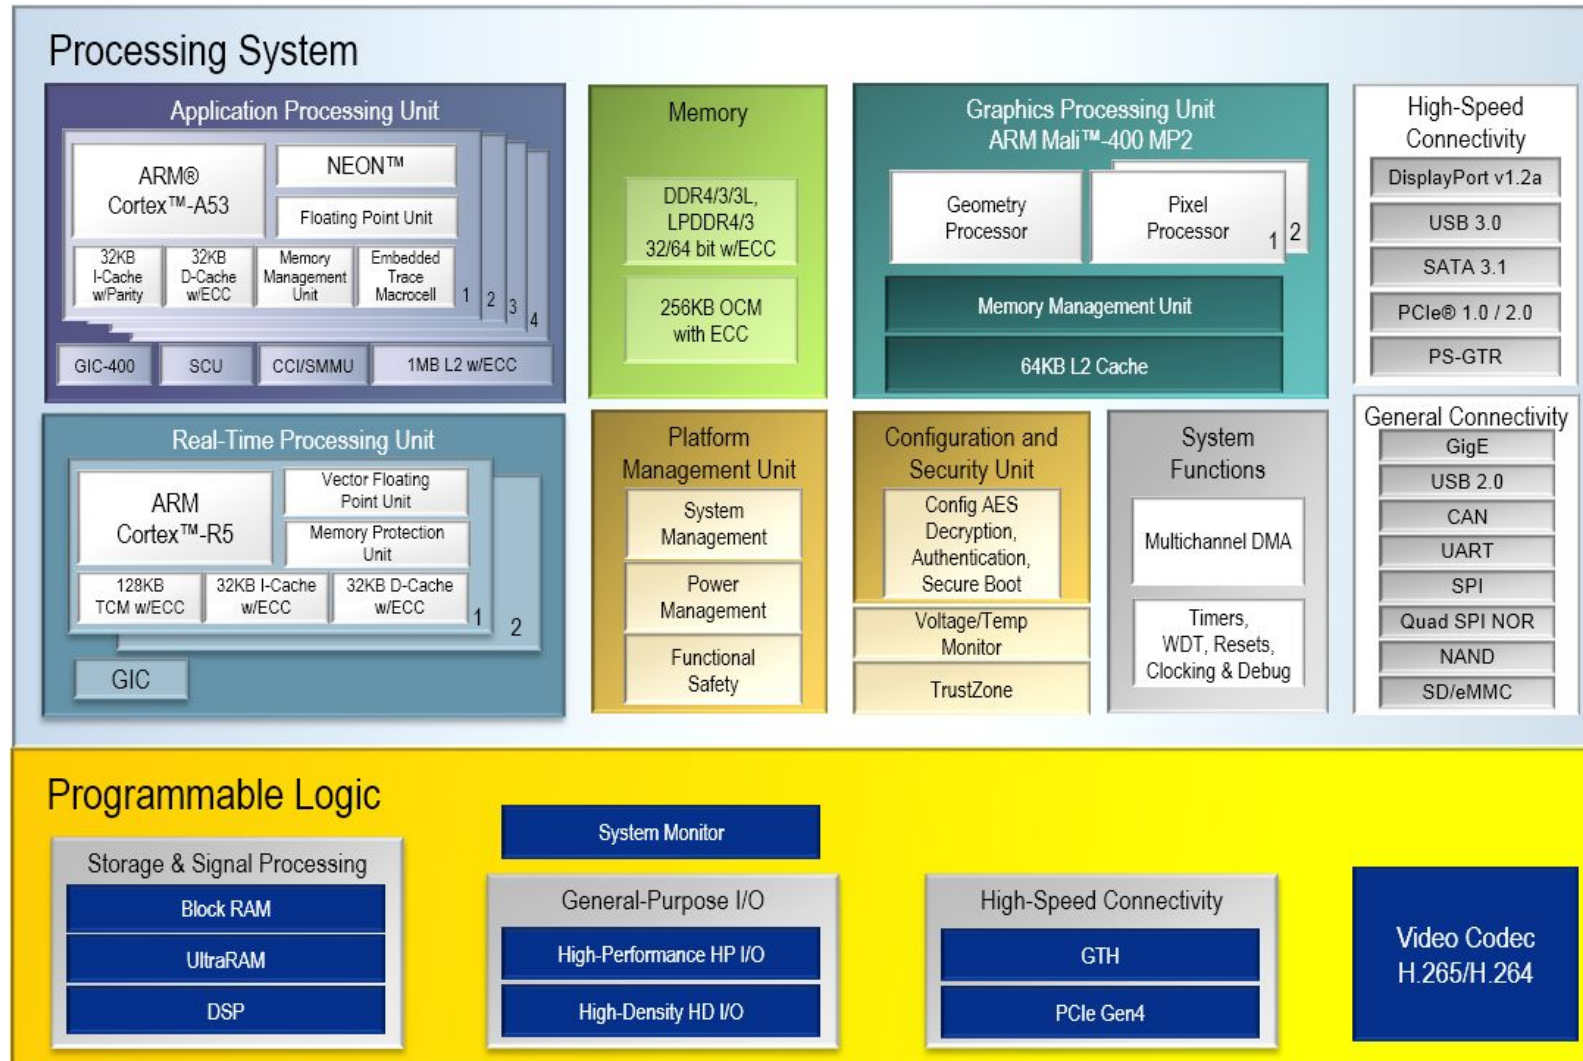
Where do I  
start with  
SDSoC?



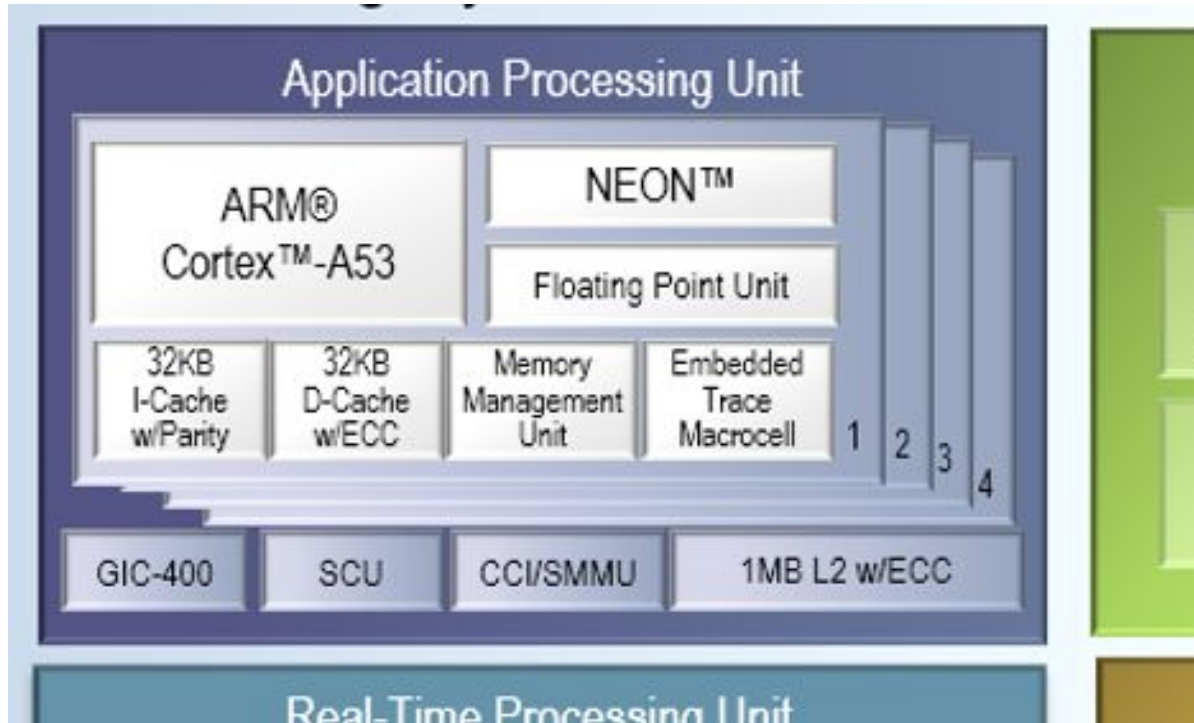
# What is Zynq?



# What is Zynq UltraScale+?

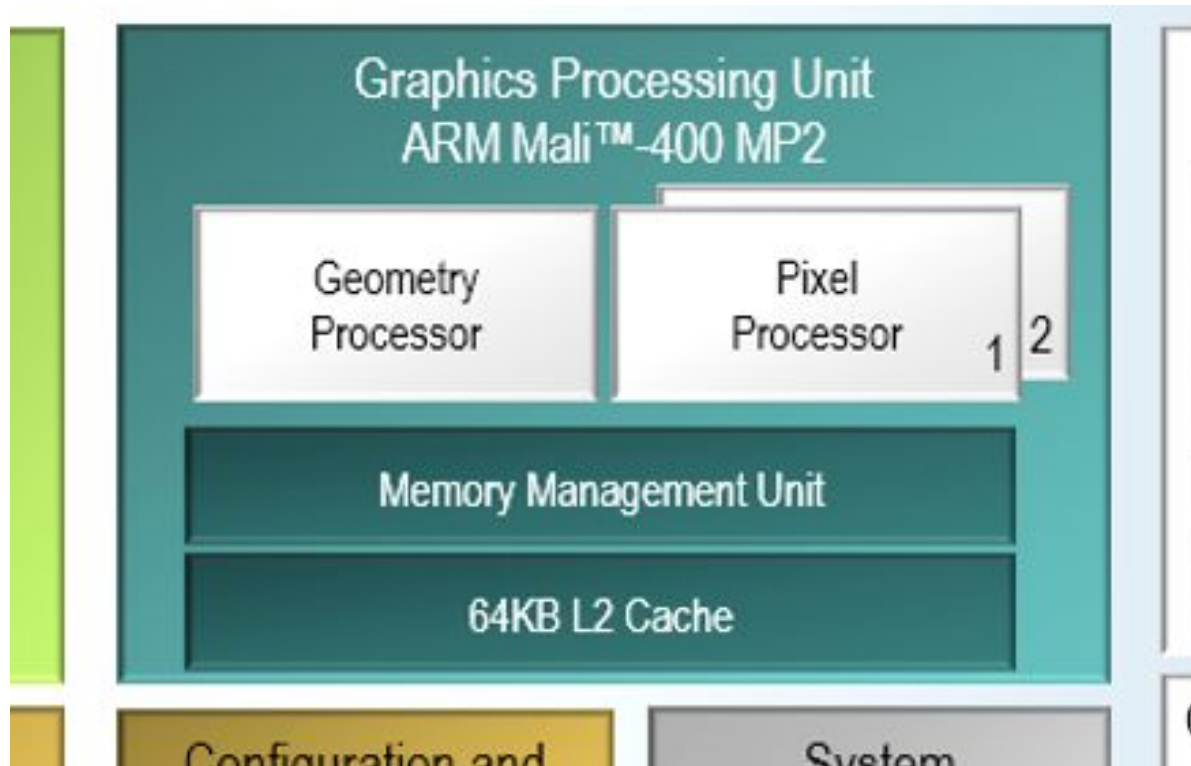


# What is Zynq UltraScale+?



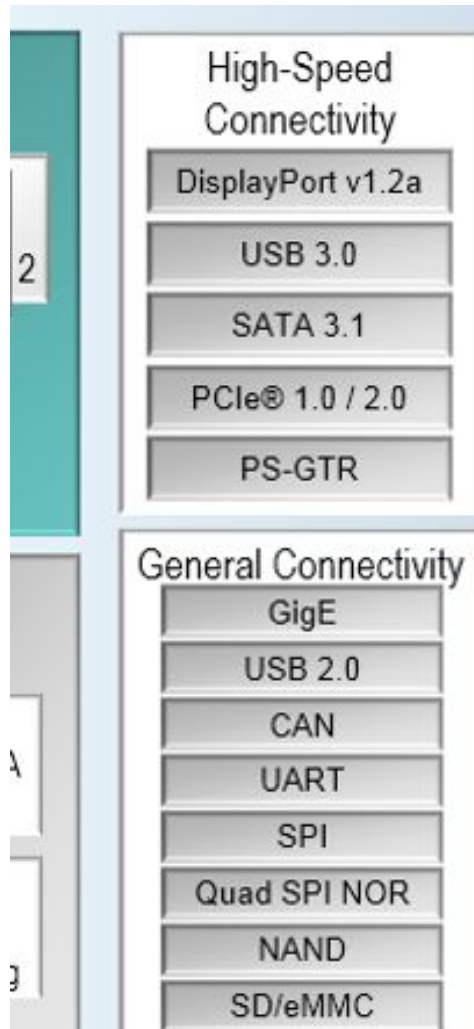
- Quad Core Arm A53
- NEON SIMD Accelerators

# What is Zynq UltraScale+?



- Quad Core Arm A53
- NEON SIMD Accelerators
- Mali GPU

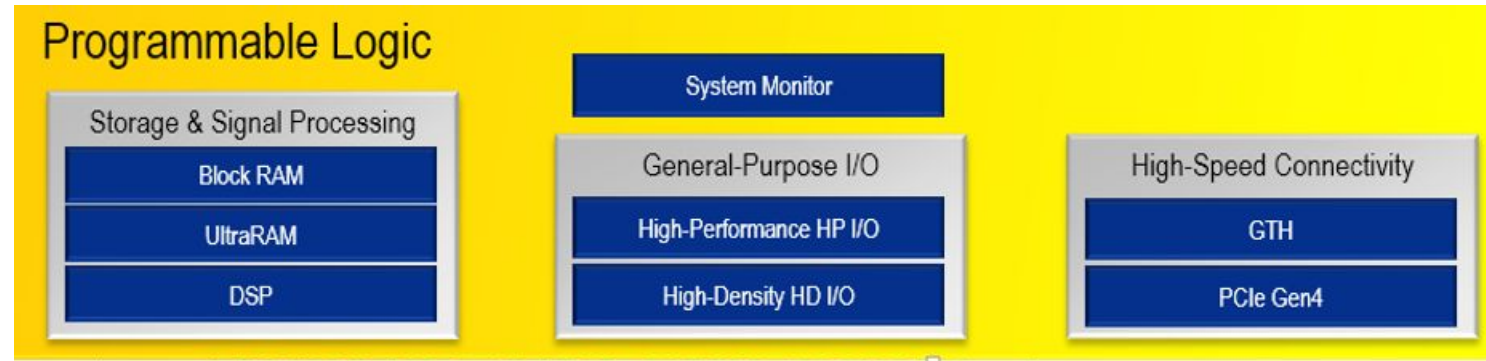
# What is Zynq UltraScale+?



- Quad Core Arm A53
- NEON SIMD Accelerators
- Mali GPU
- Hardened IP
- VCU (Video Codec Unit)

# What is Zynq UltraScale+?



- Quad Core Arm A53
- NEON SIMD Accelerators
- Mali GPU
- Hardened IP
- VCU (Video Codec Unit)
- Variable amounts of Programmable Logic (amount of acceleration possible)
- Variable grades of Programmable Logic (speed)
- Variable fabric families of Programmable Logic (performance)





# What is Zynq UltraScale+?

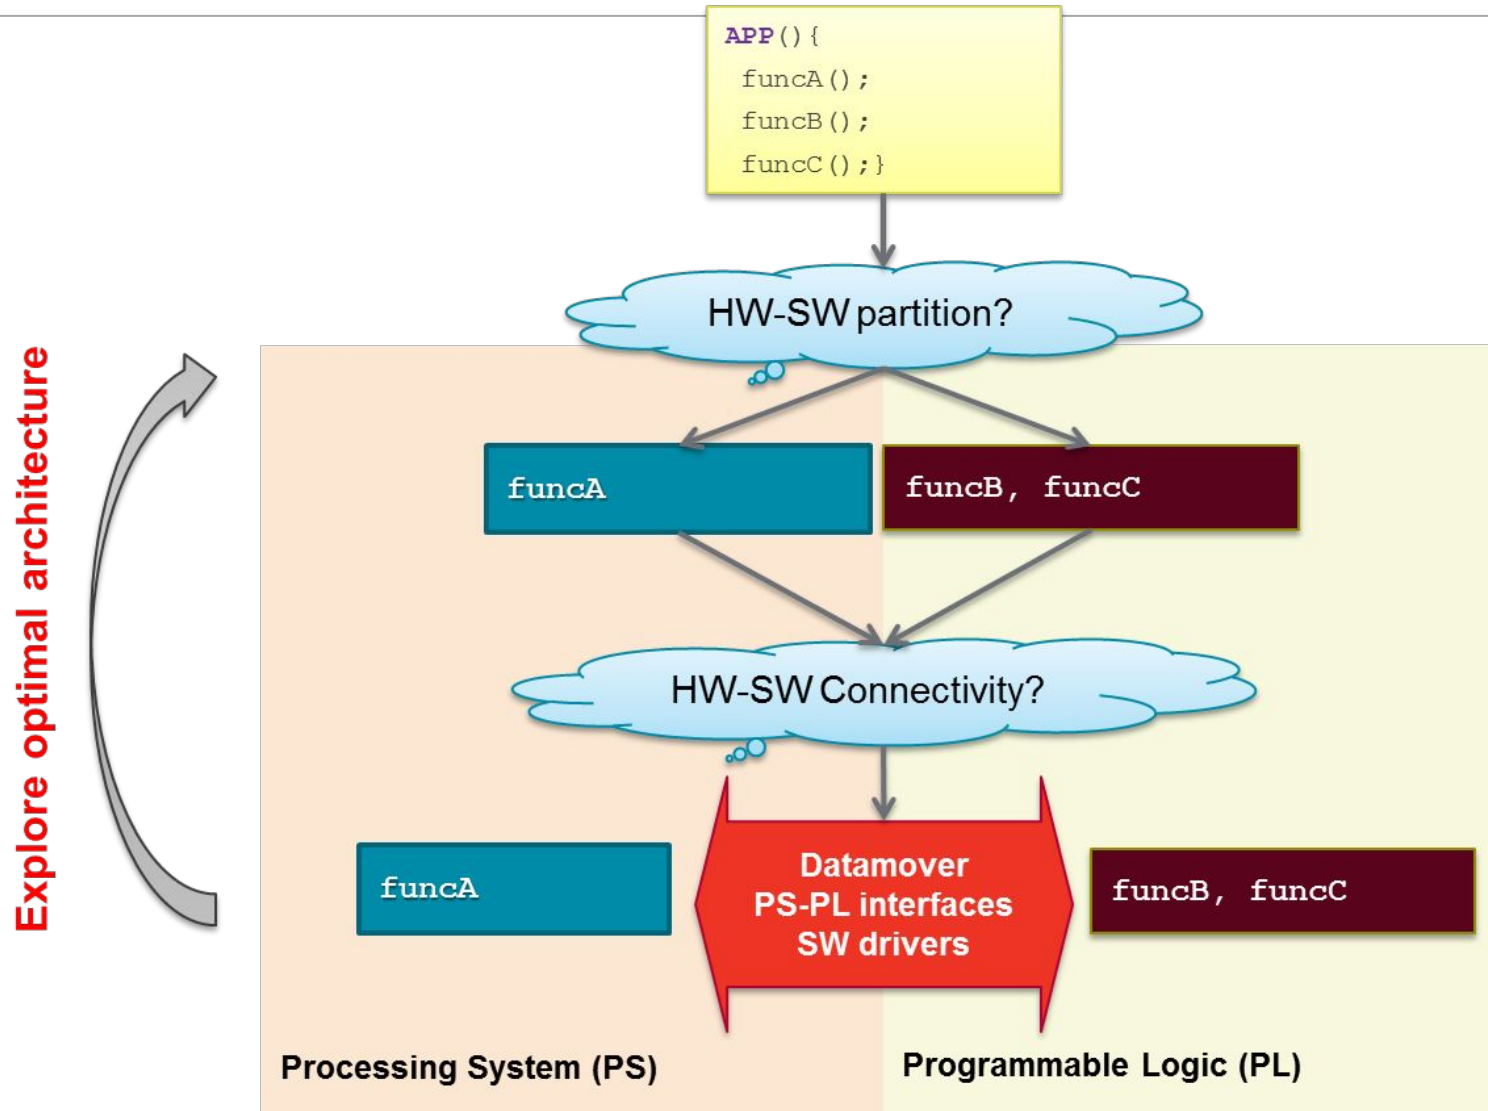
- EX: Ultra96 \$249; comes with SDSoC License
- 96Boards Compatible



What is my  
design  
flow?

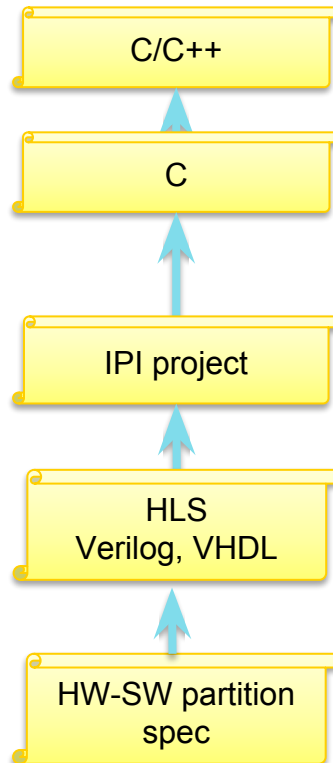


# Typical Zynq Development Flow




# Before SDSoC:

- Involves Multiple Disciplines
- Direct manual coordination for solution

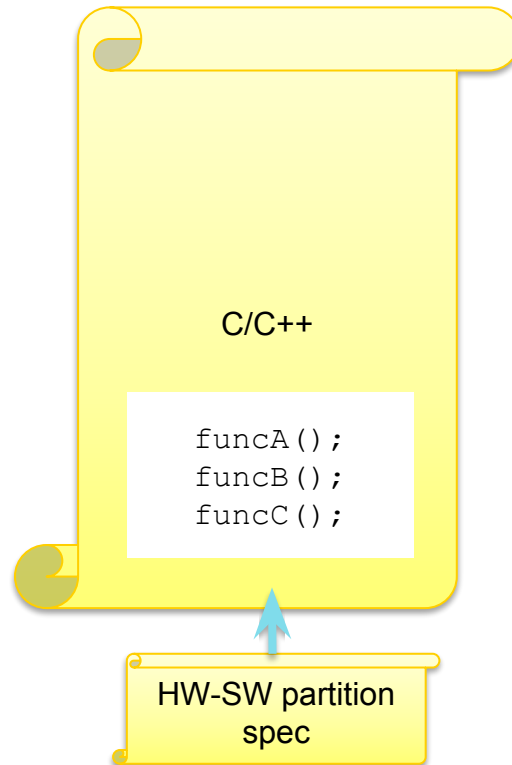


# After SDSoC:

- Remove the manual design of SW drivers and HW connectivity

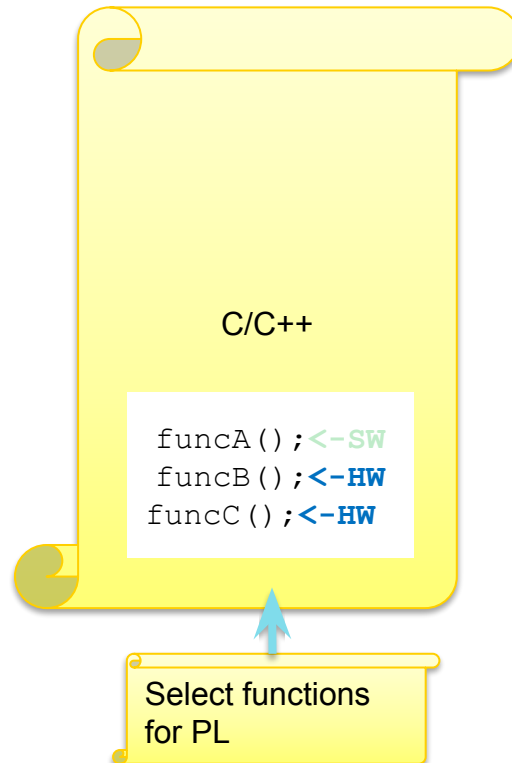


# After SDSoc:



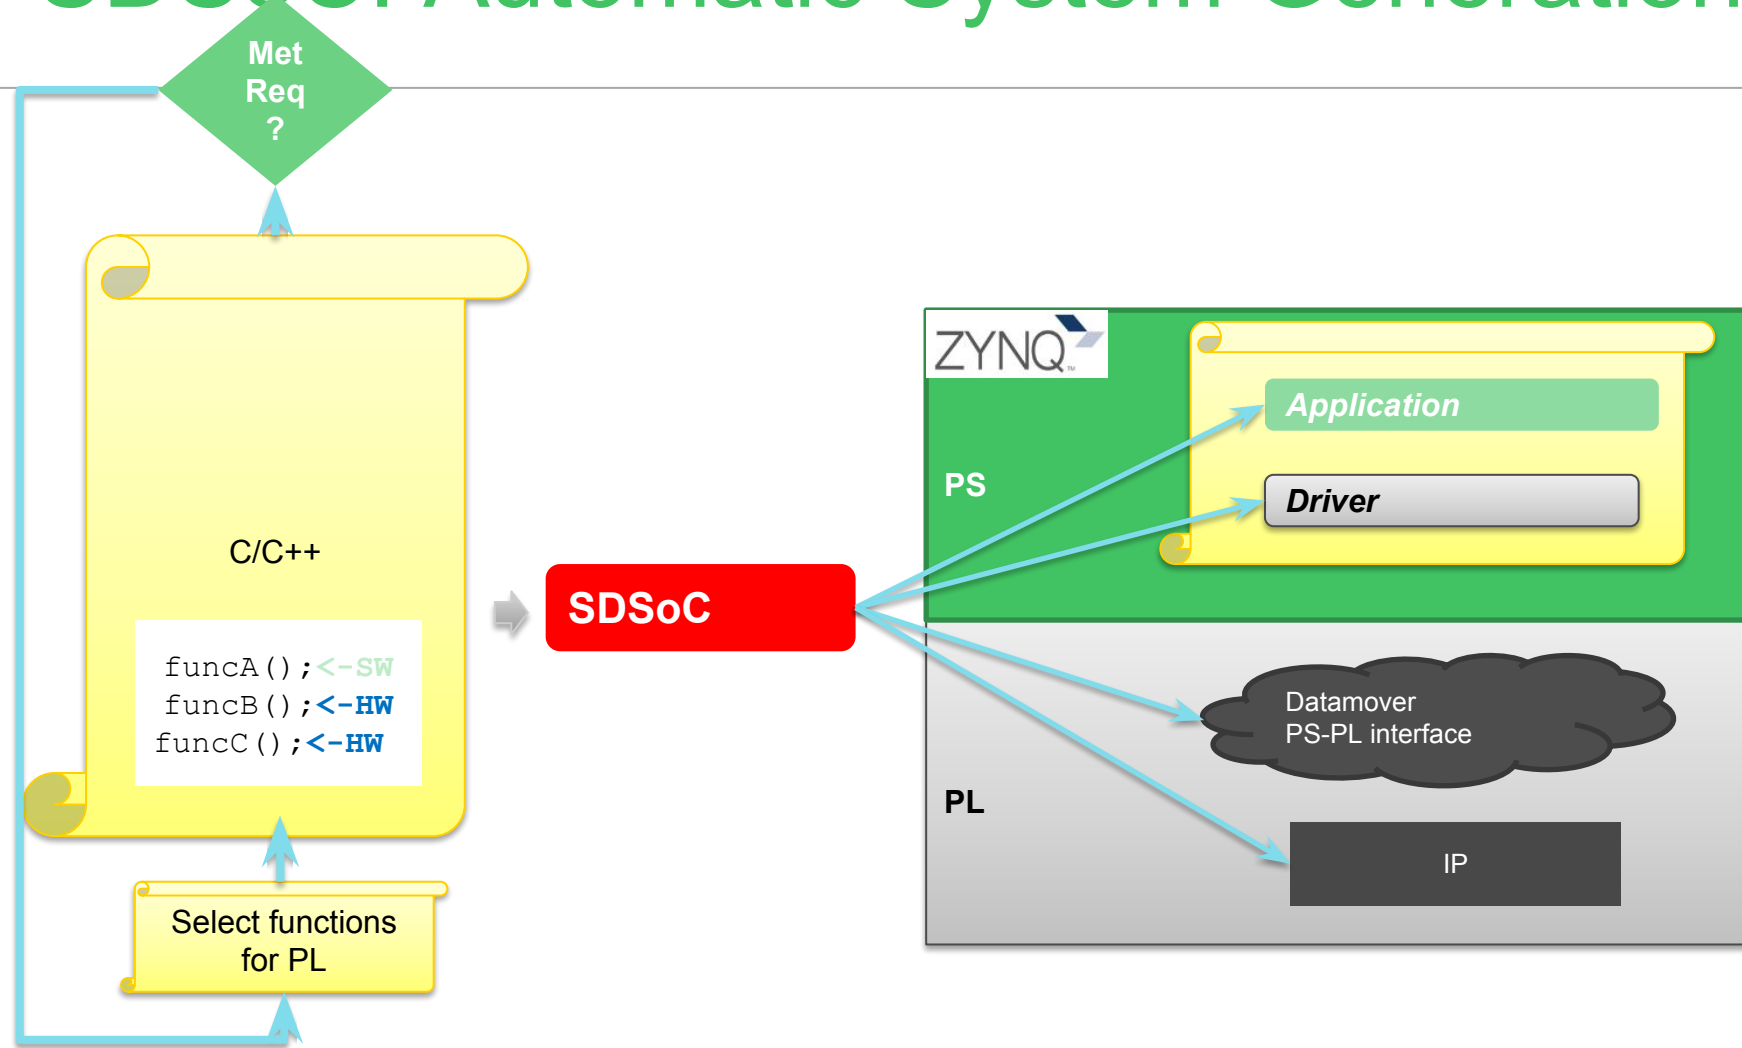
- Remove the manual design of SW drivers and HW connectivity
- Use the C/C++ end application as the input calling the user algorithm IPs as function calls

# After SDSoc:



- Remove the manual design of SW drivers and HW connectivity
- Use the C/C++ end application as the input calling the user algorithm IPs as function calls
- Partition set of functions to Programmable Logic by a single click

# After SDSoC: Automatic System Generation



Platforms...  
what are  
they?

# Platforms...what are they?

---

- A *platform* is a base system designed for reuse, our sandbox
  - Built using Vivado, SDK (now optional) and OS tools
  - Your Hardware is Defined: Processing system, I/O subsystems, memory interfaces, ..., with a well-defined port interface (AXI, AXI-S, clock, reset, interrupt)
  - Software Definition: OS, device drivers, boot loaders, file system, libraries,...




# Platforms...what are they?

---

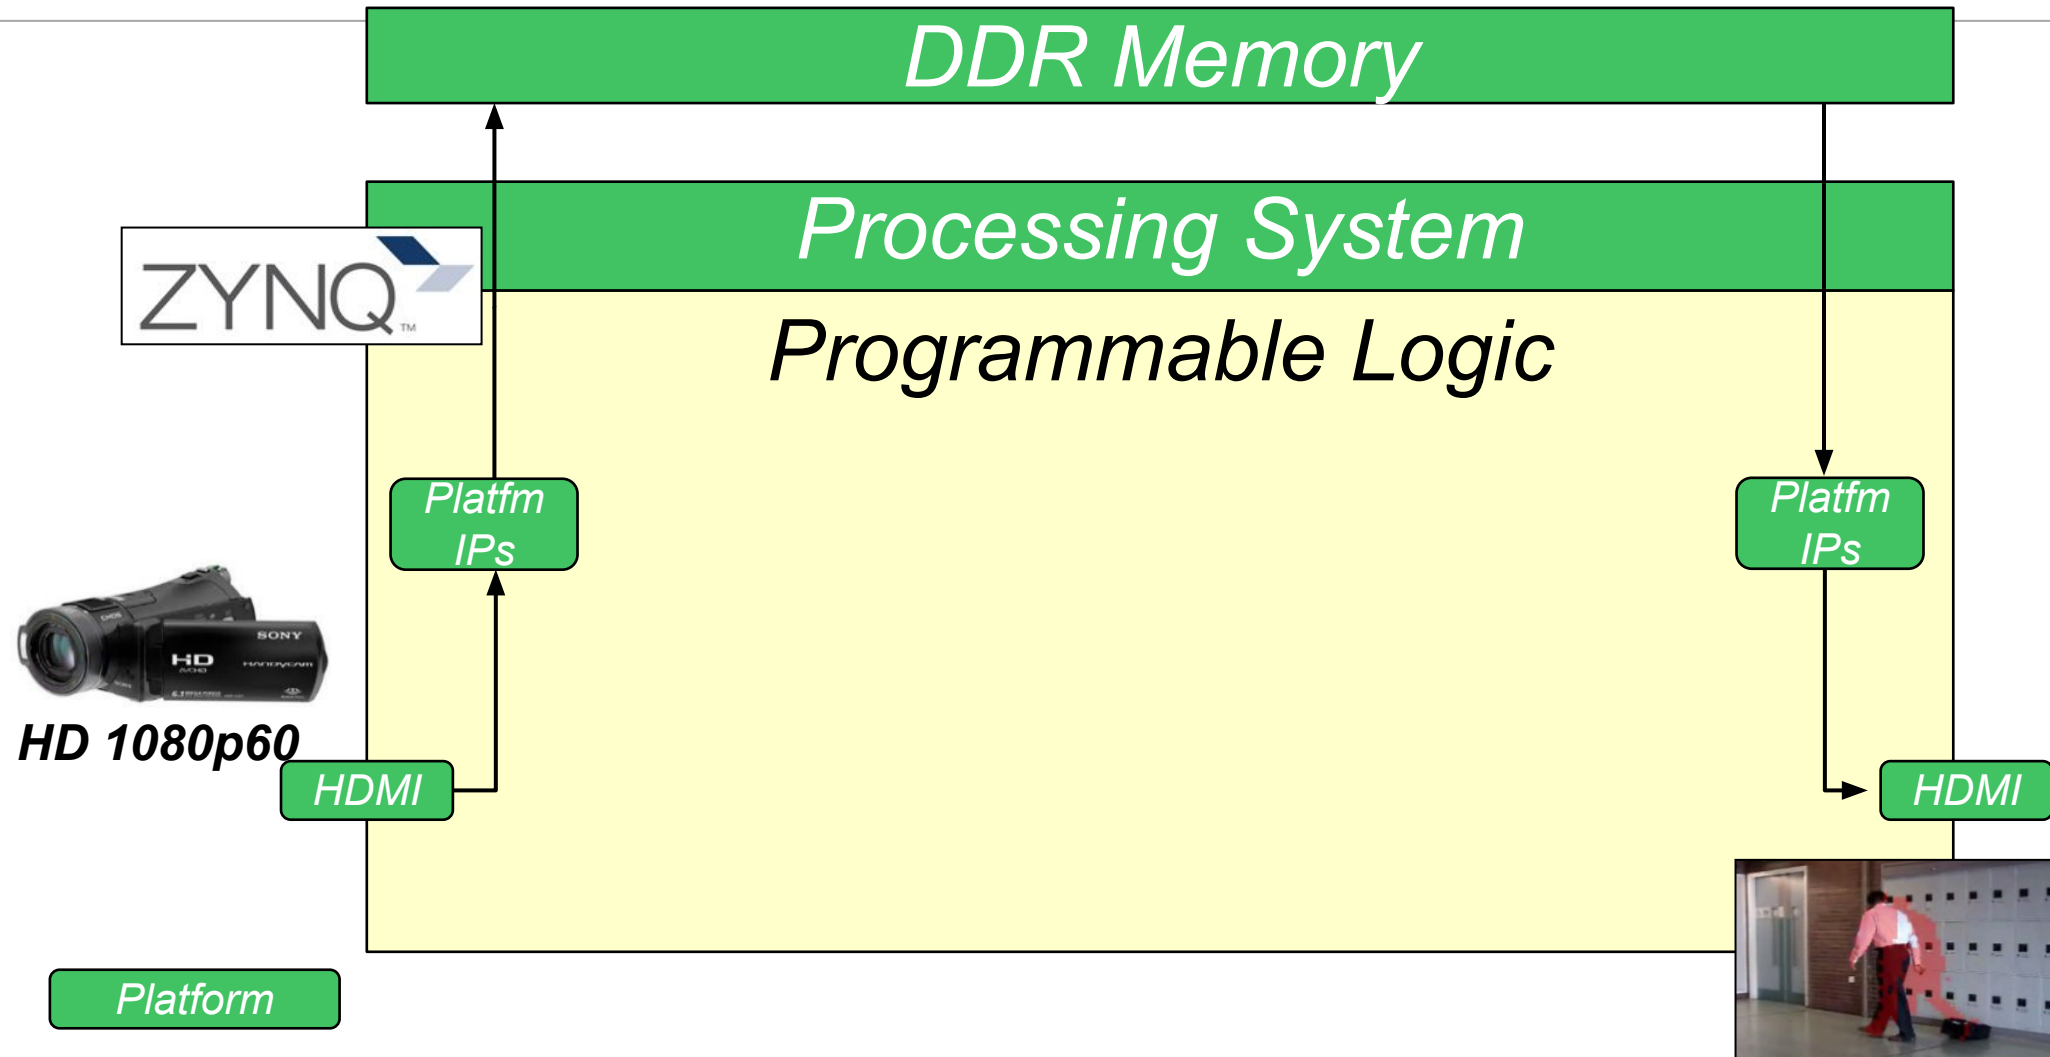
- SDSoC extends the platform with application-specific hardware and software
  - Complete solution generated from C/C++ source
  - Solution can include RTL packaged
- SDSoC treats platforms as independent solution spaces
  - Generate IPs for your solution based on the resources available
  - Each solution is tailored for each platform
  - A solution can have the target platform switched through little more than a few GUI clicks!
- Platforms must include
  - Processing IP Block
  - At least one general purpose AXI master port

# XPFM: A Pictorial

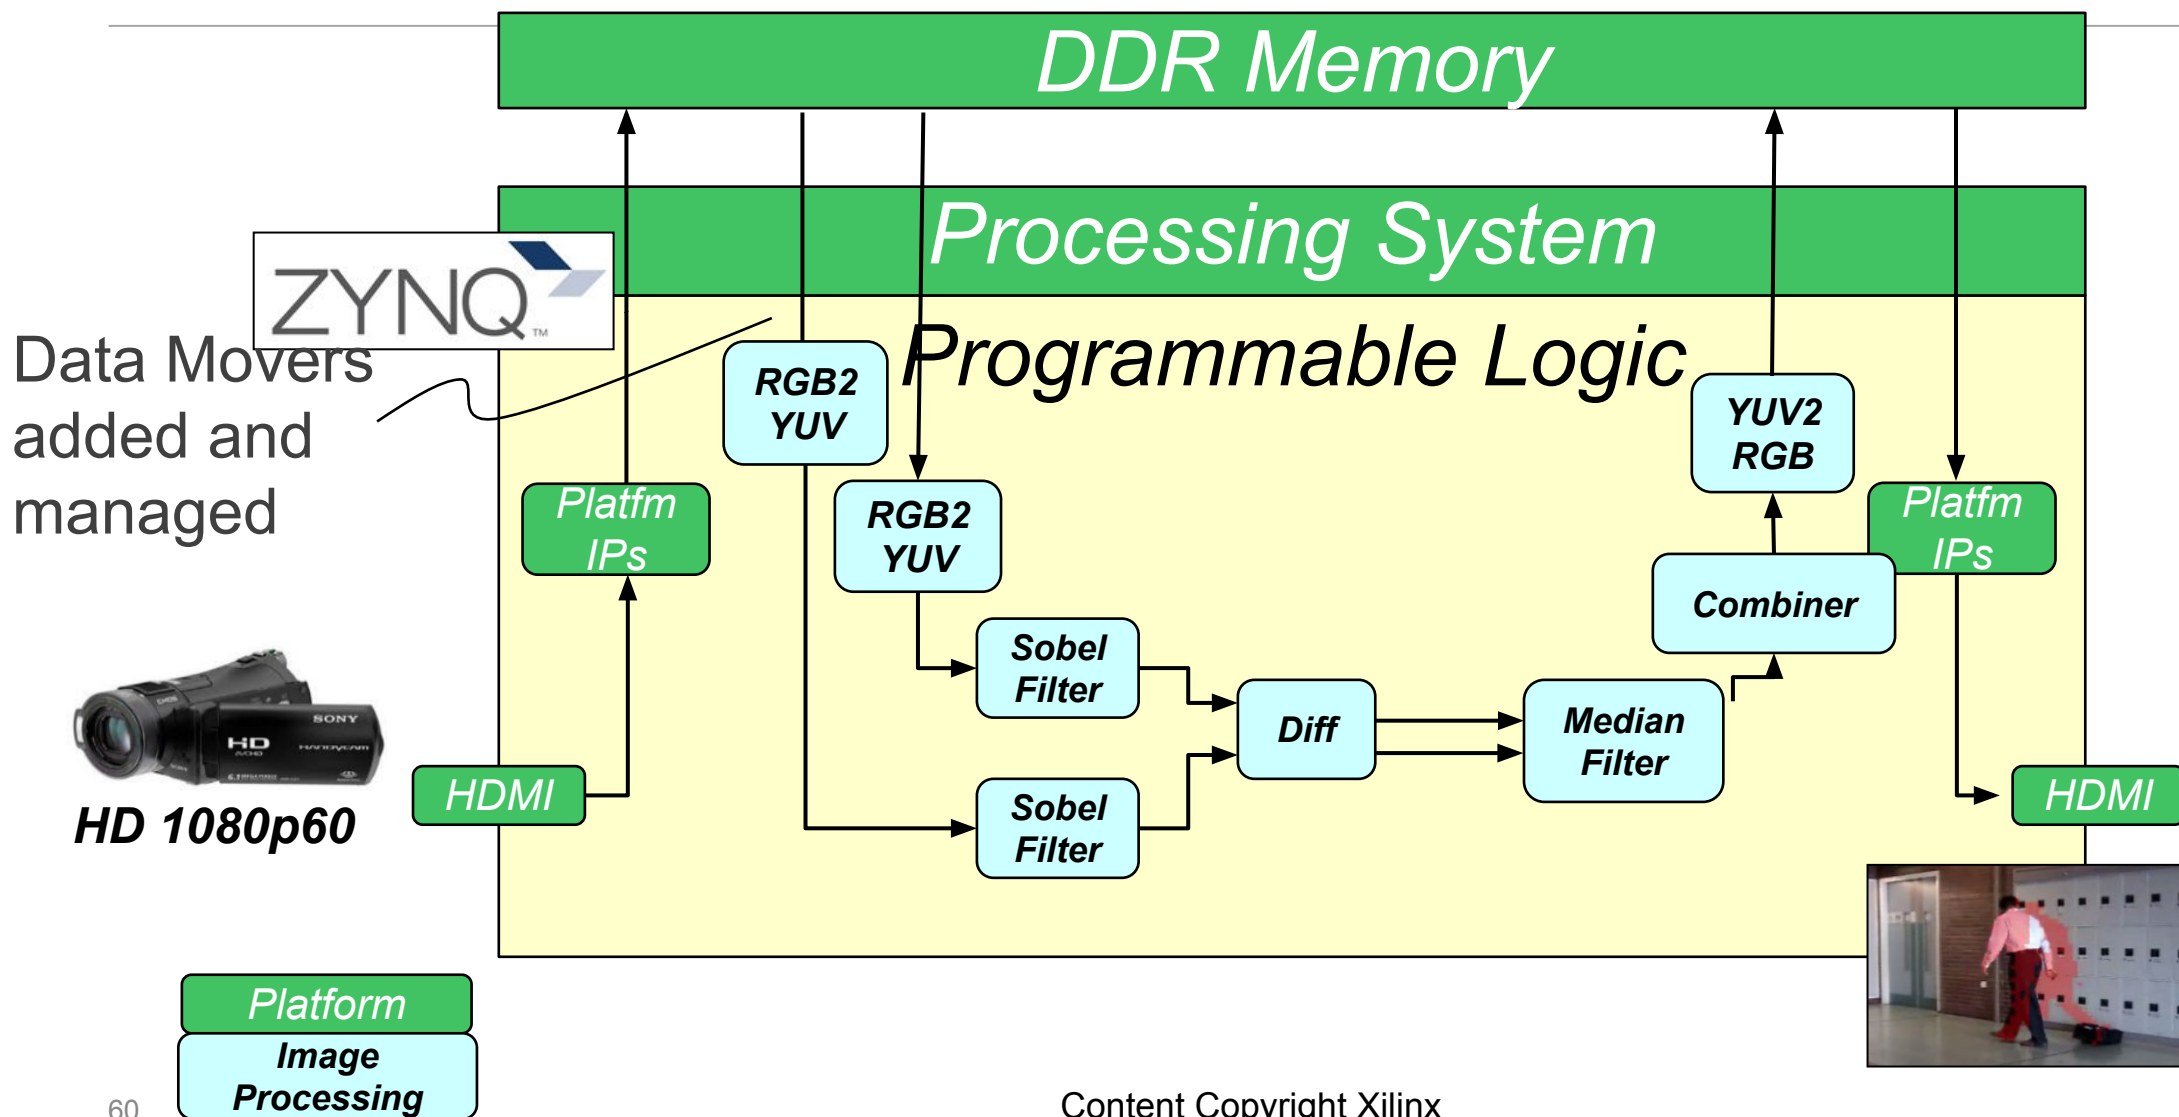
- XPFM
  - Platform Reference to the location of the DSA and SPFM
- BD
  - Block Design
- DSA
  - Device Support Archive
- SPFM
  - Software Platform Descriptor XML
- Black Box
  - SDSoc created IP



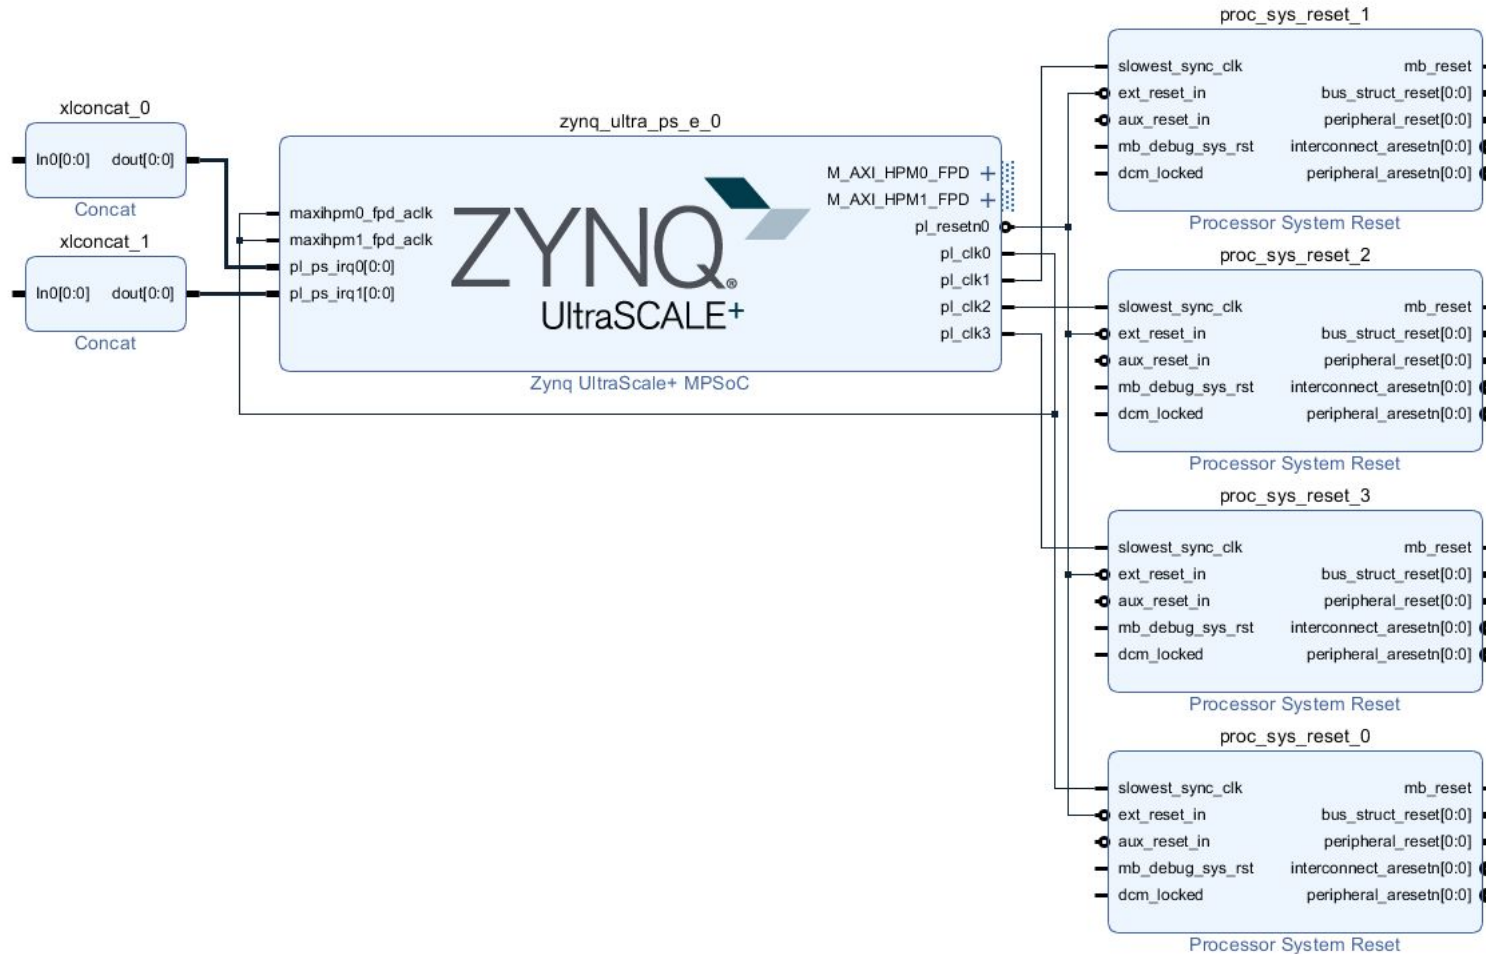
# Example: Optical Flow Design



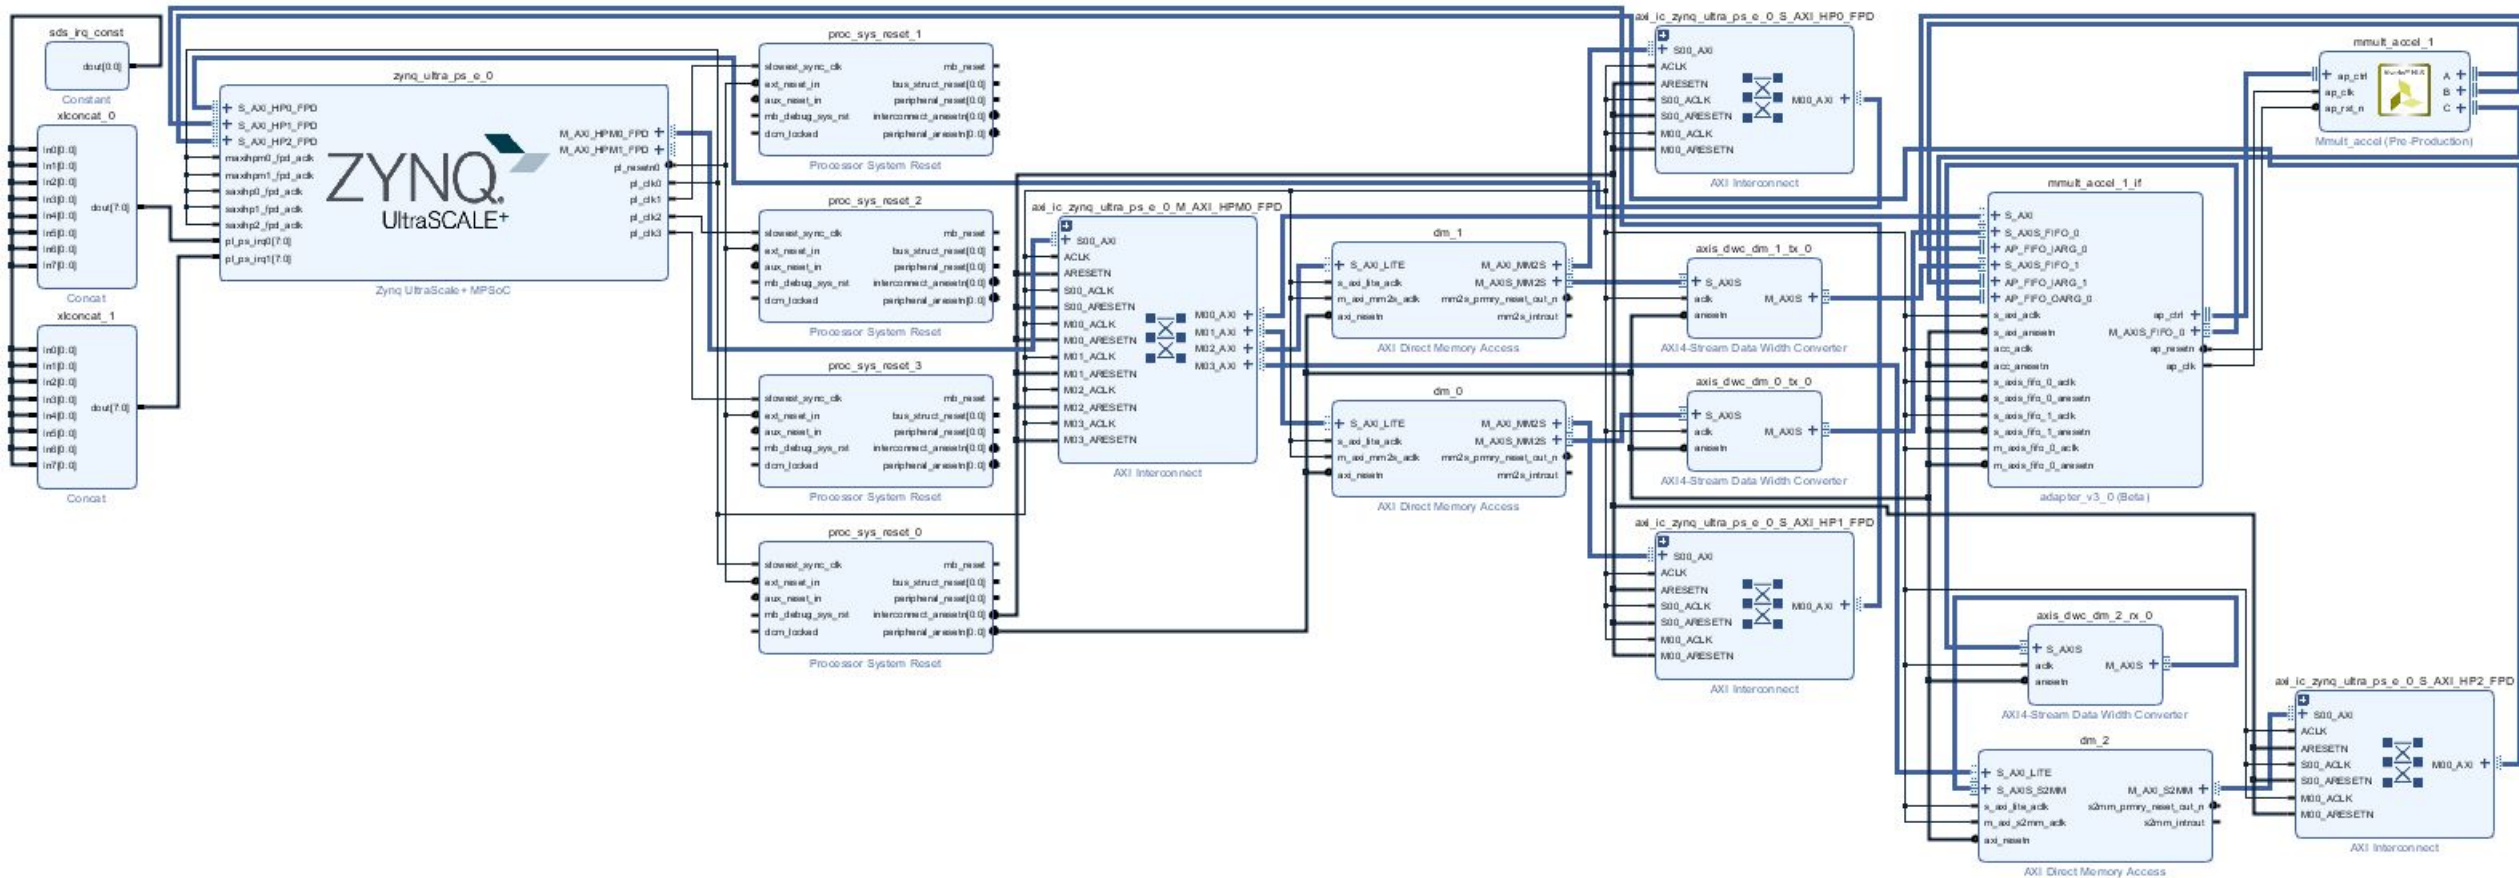
# Example: Motion Detection Application



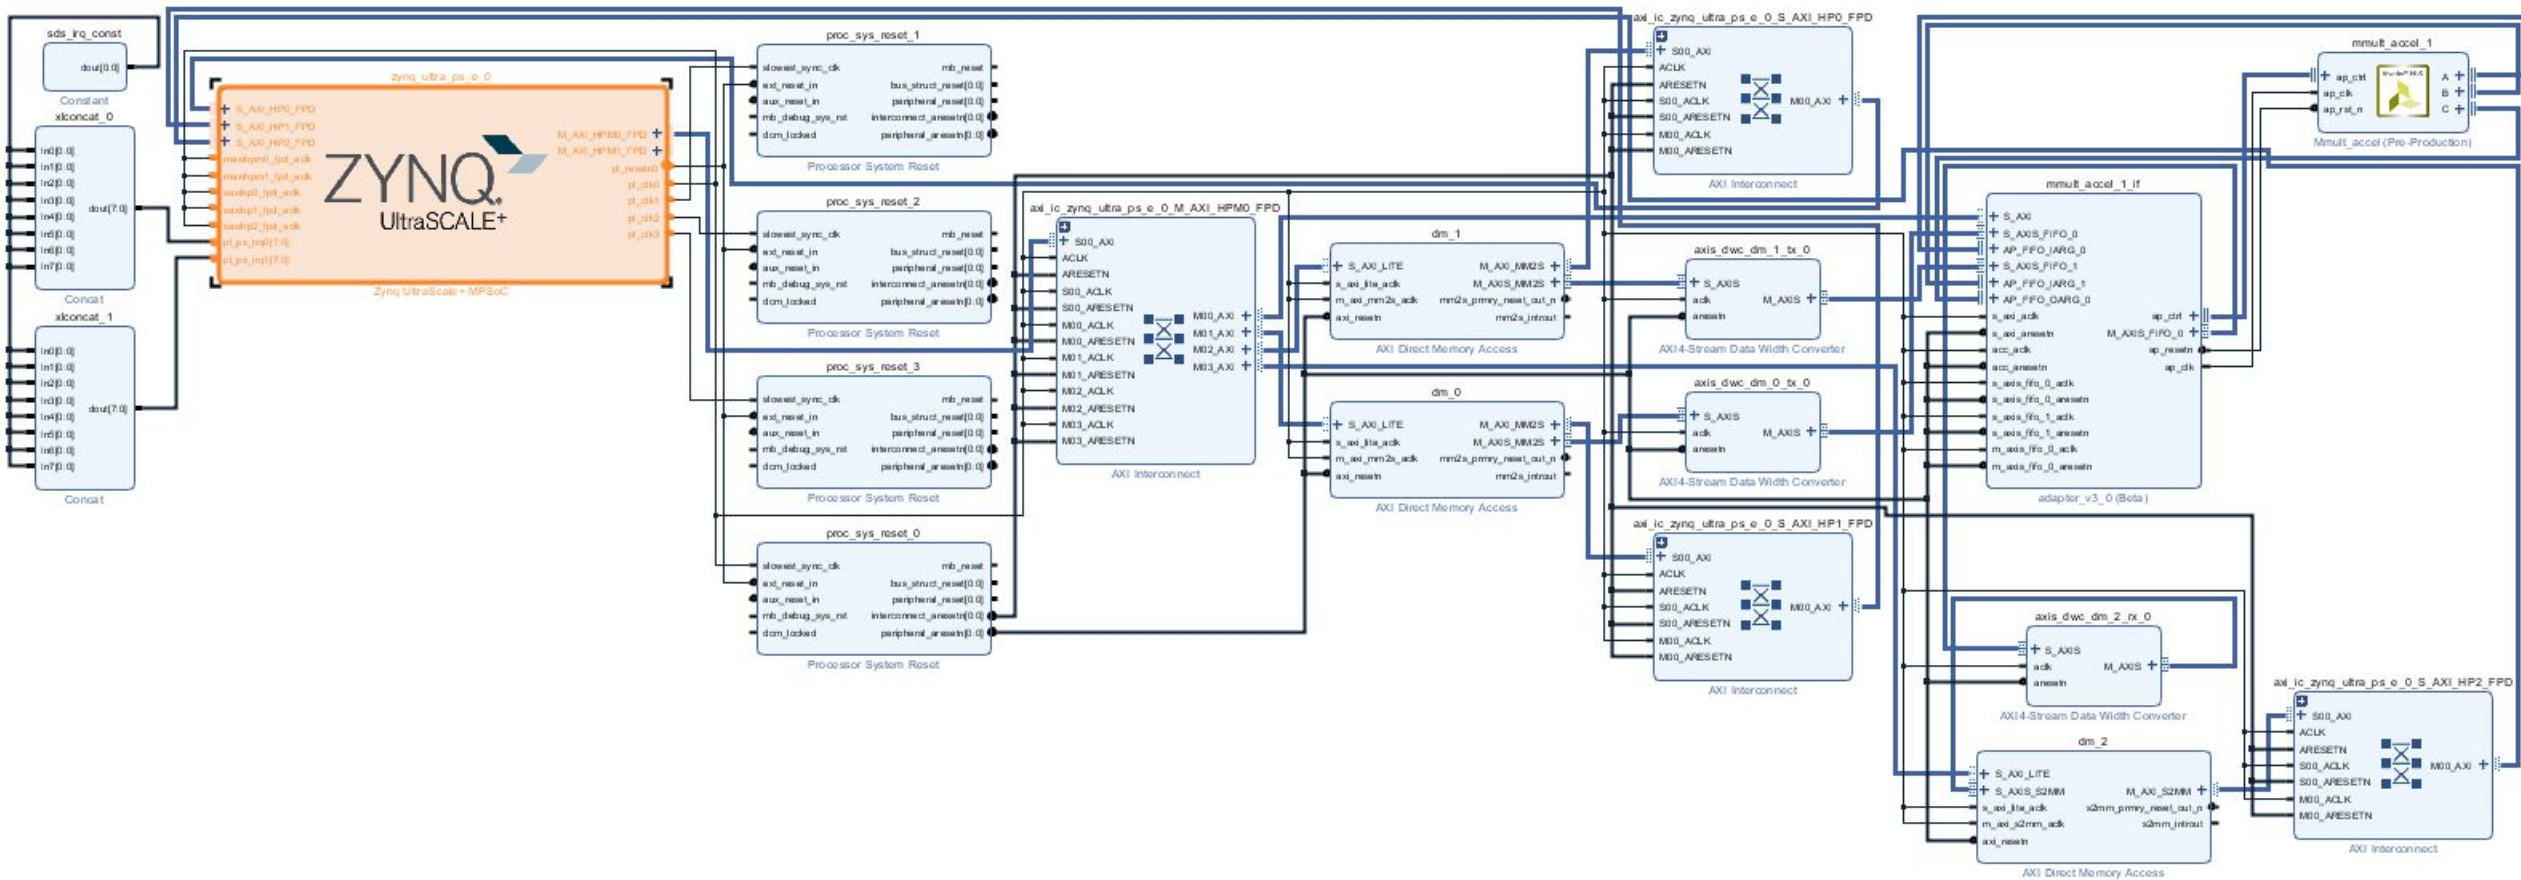
# Example: Ultra96 Platform



# Example: Ultra96 Platform with mmult

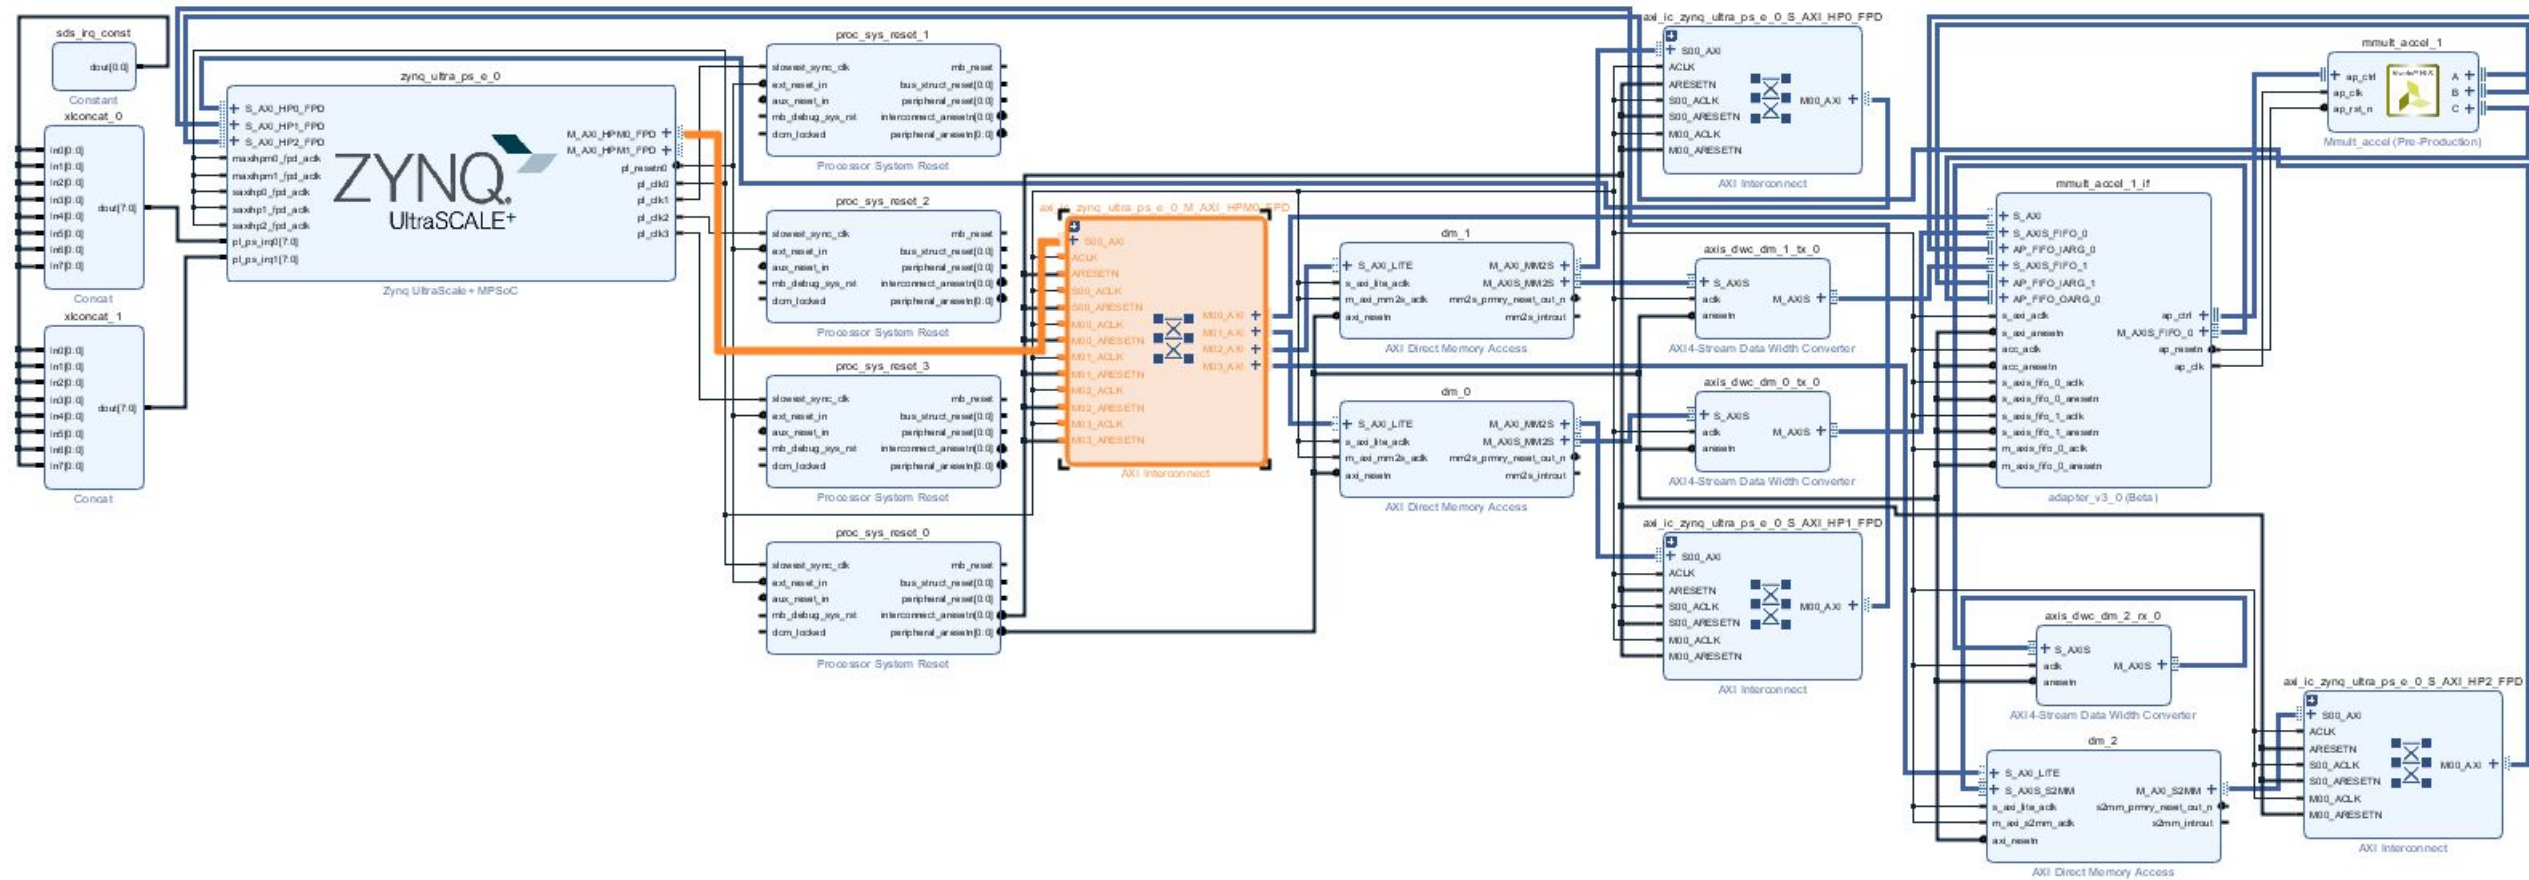


## Example: Ultra96 Platform with mmult



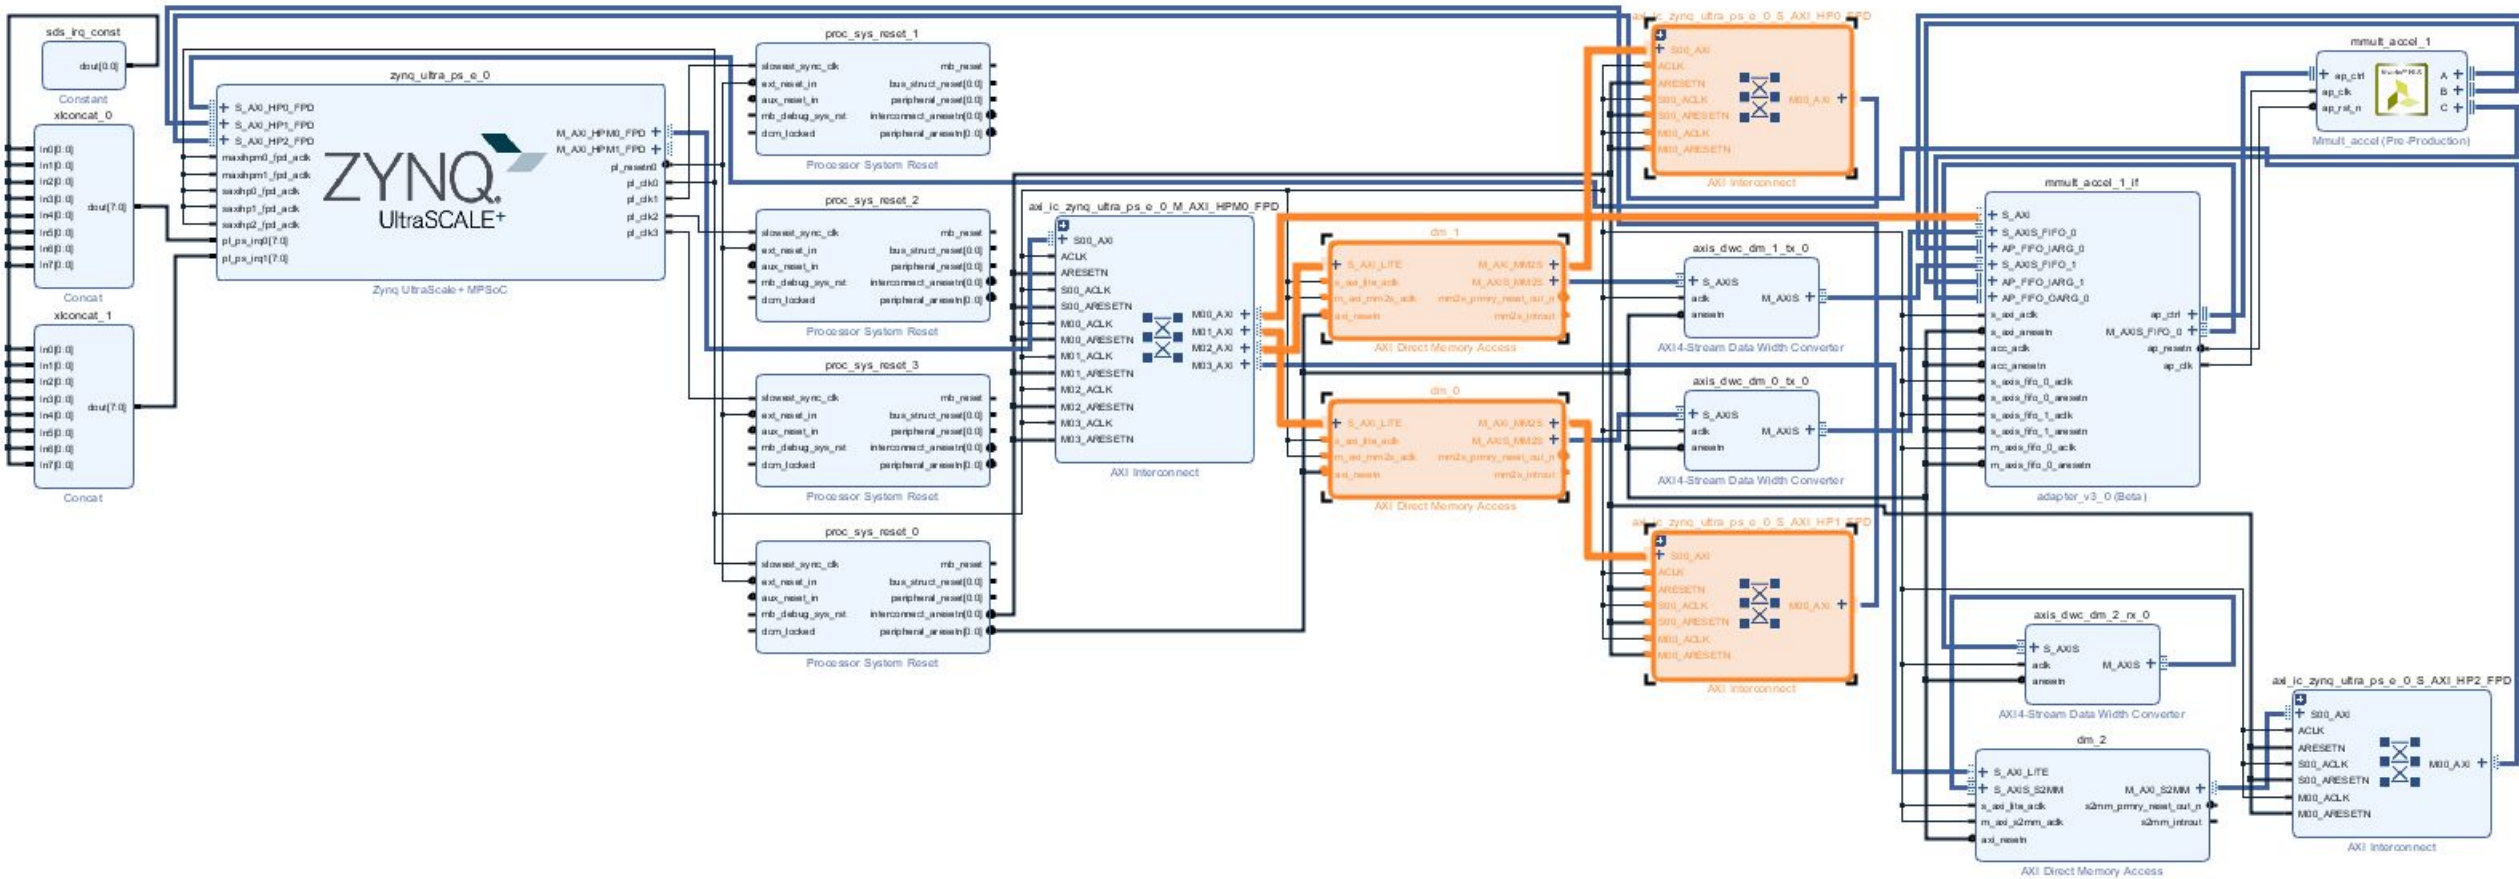


# Example: Ultra96 Platform with mmult

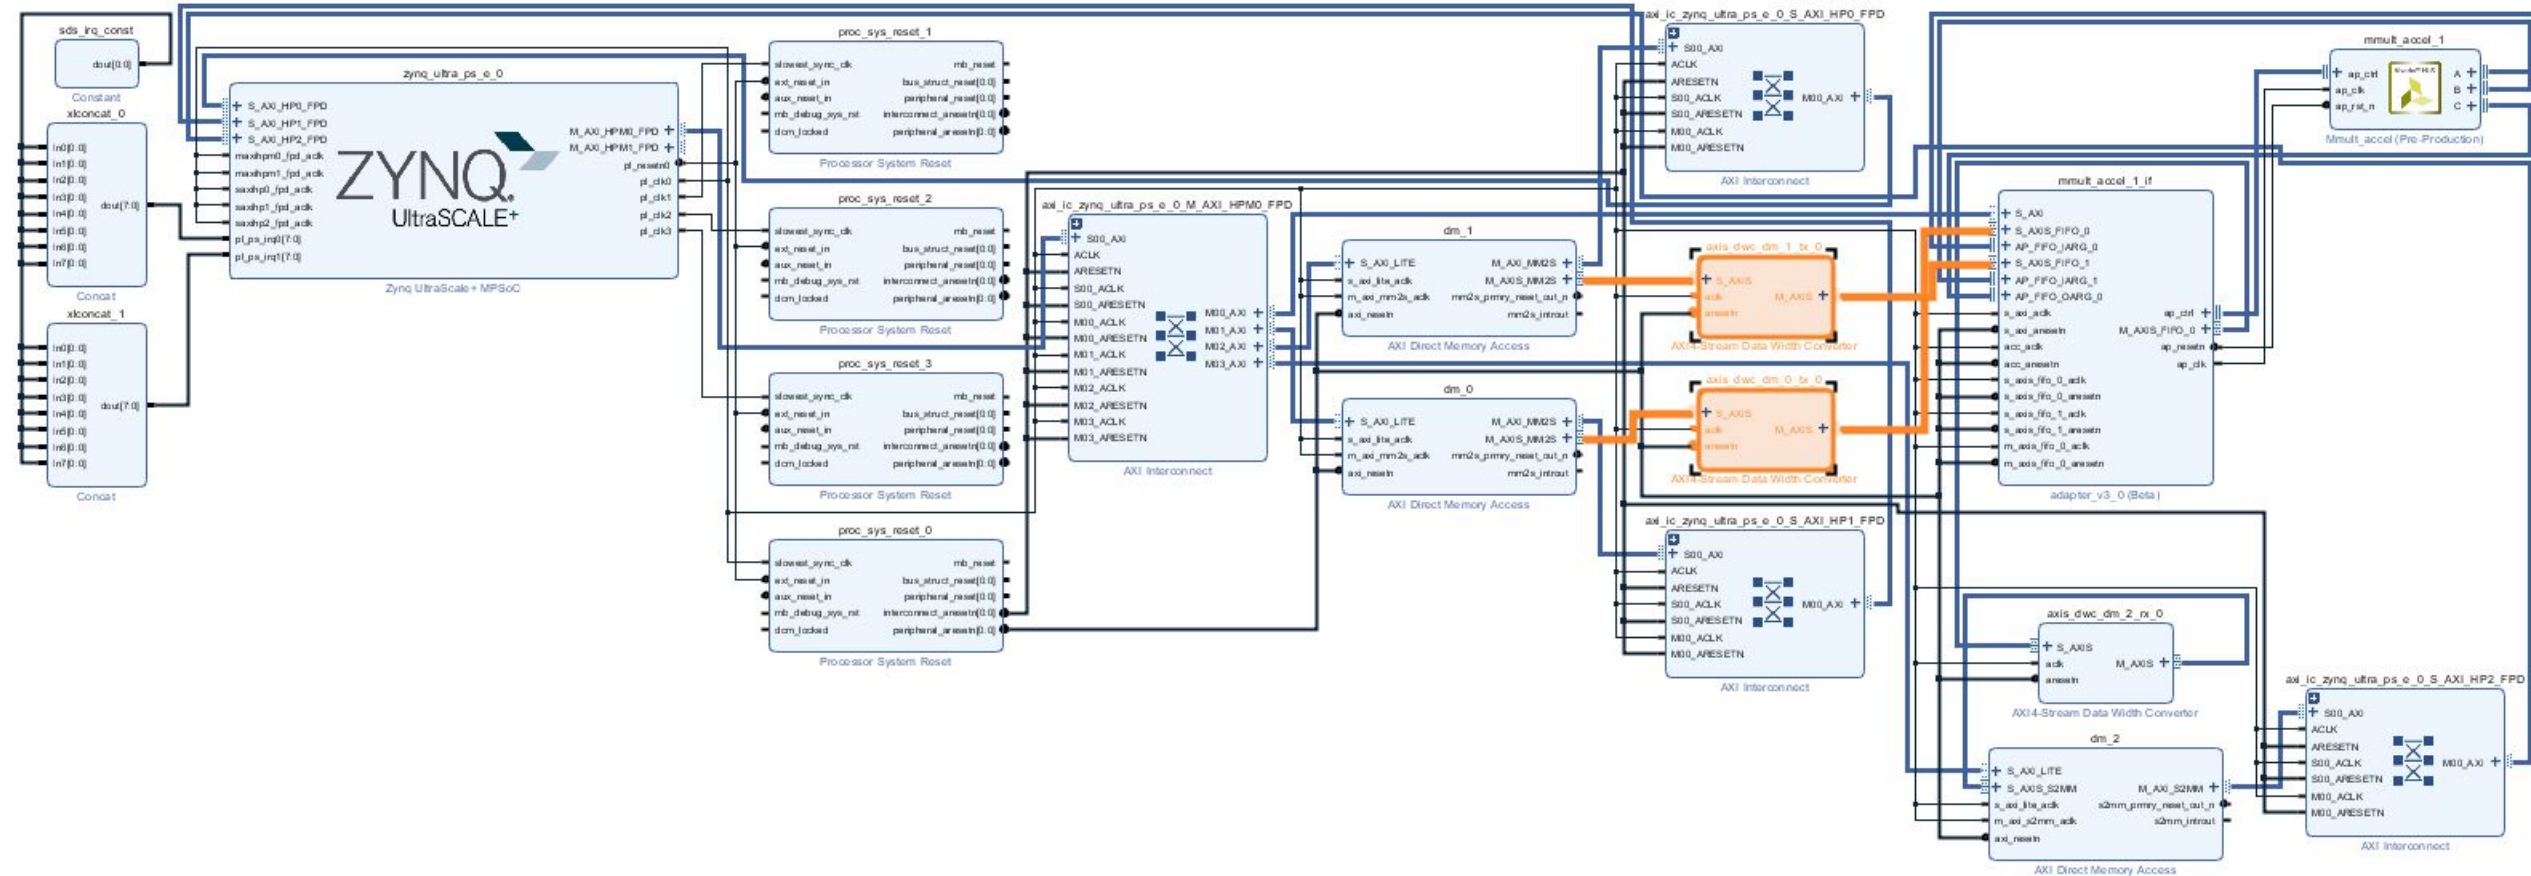




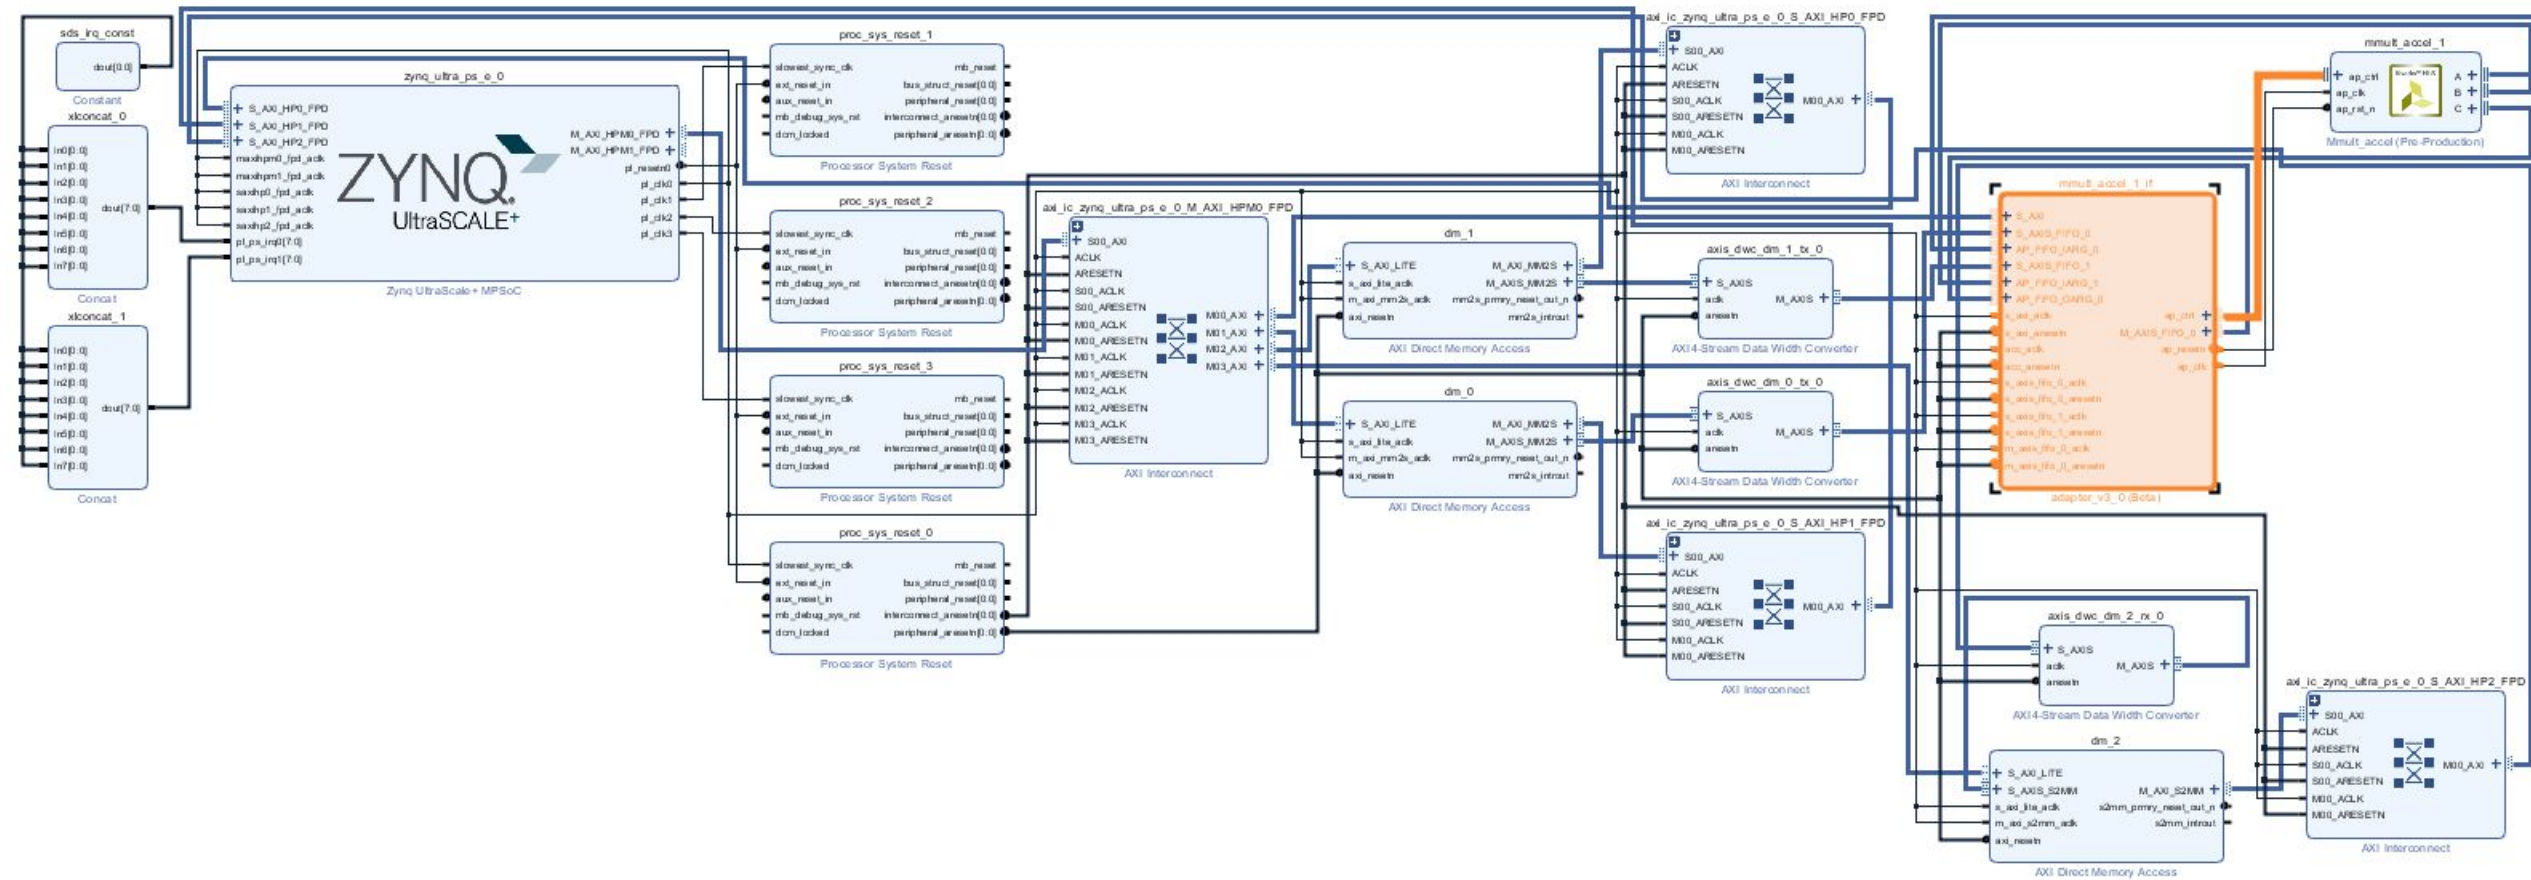
# Example: Ultra96 Platform with mmult



# Example: Ultra96 Platform with mmult

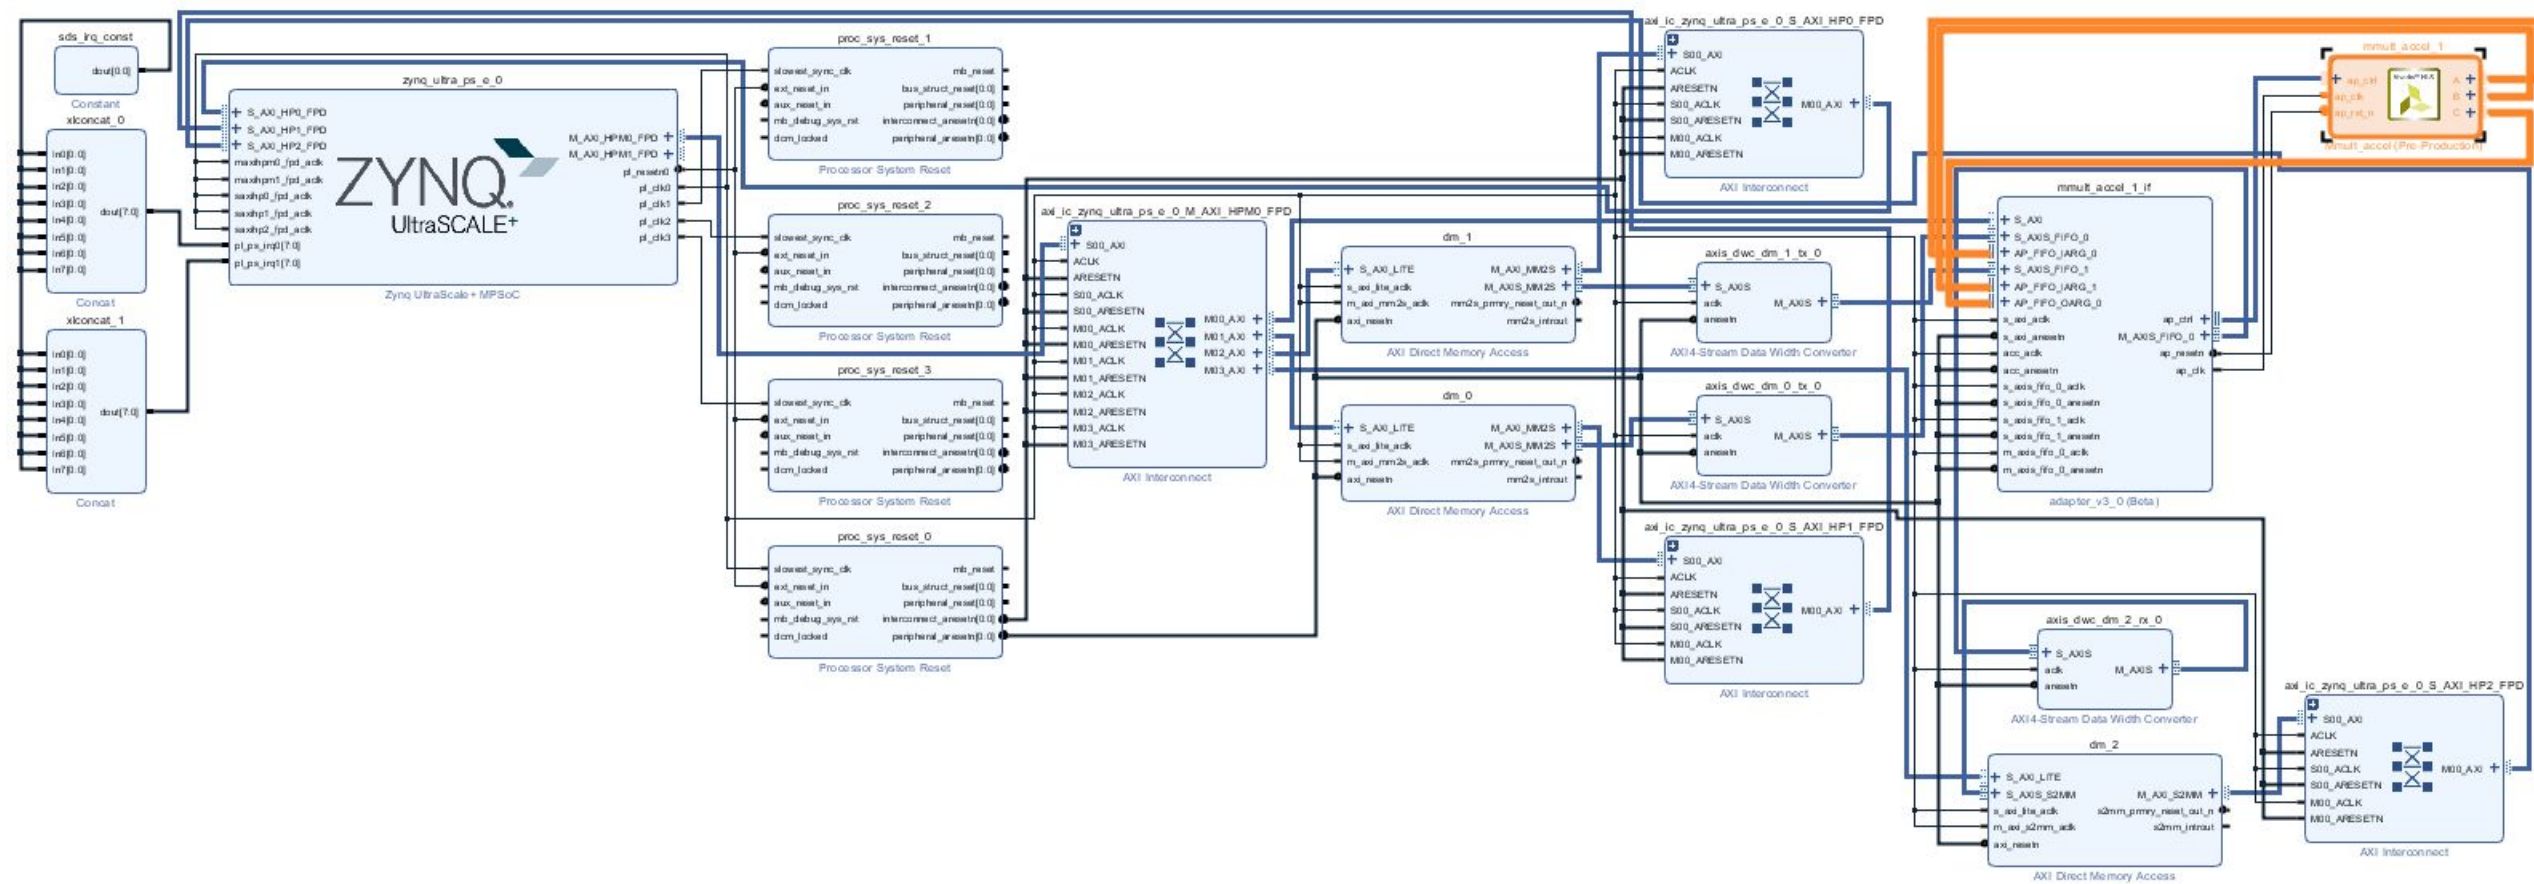


## Example: Ultra96 Platform with mmult

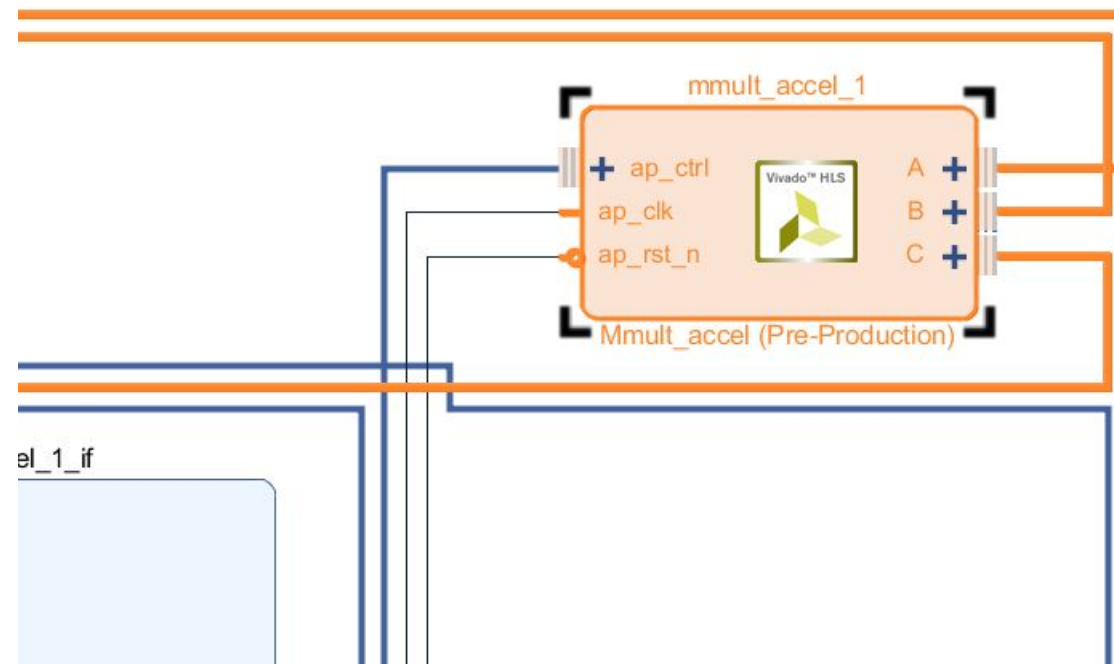




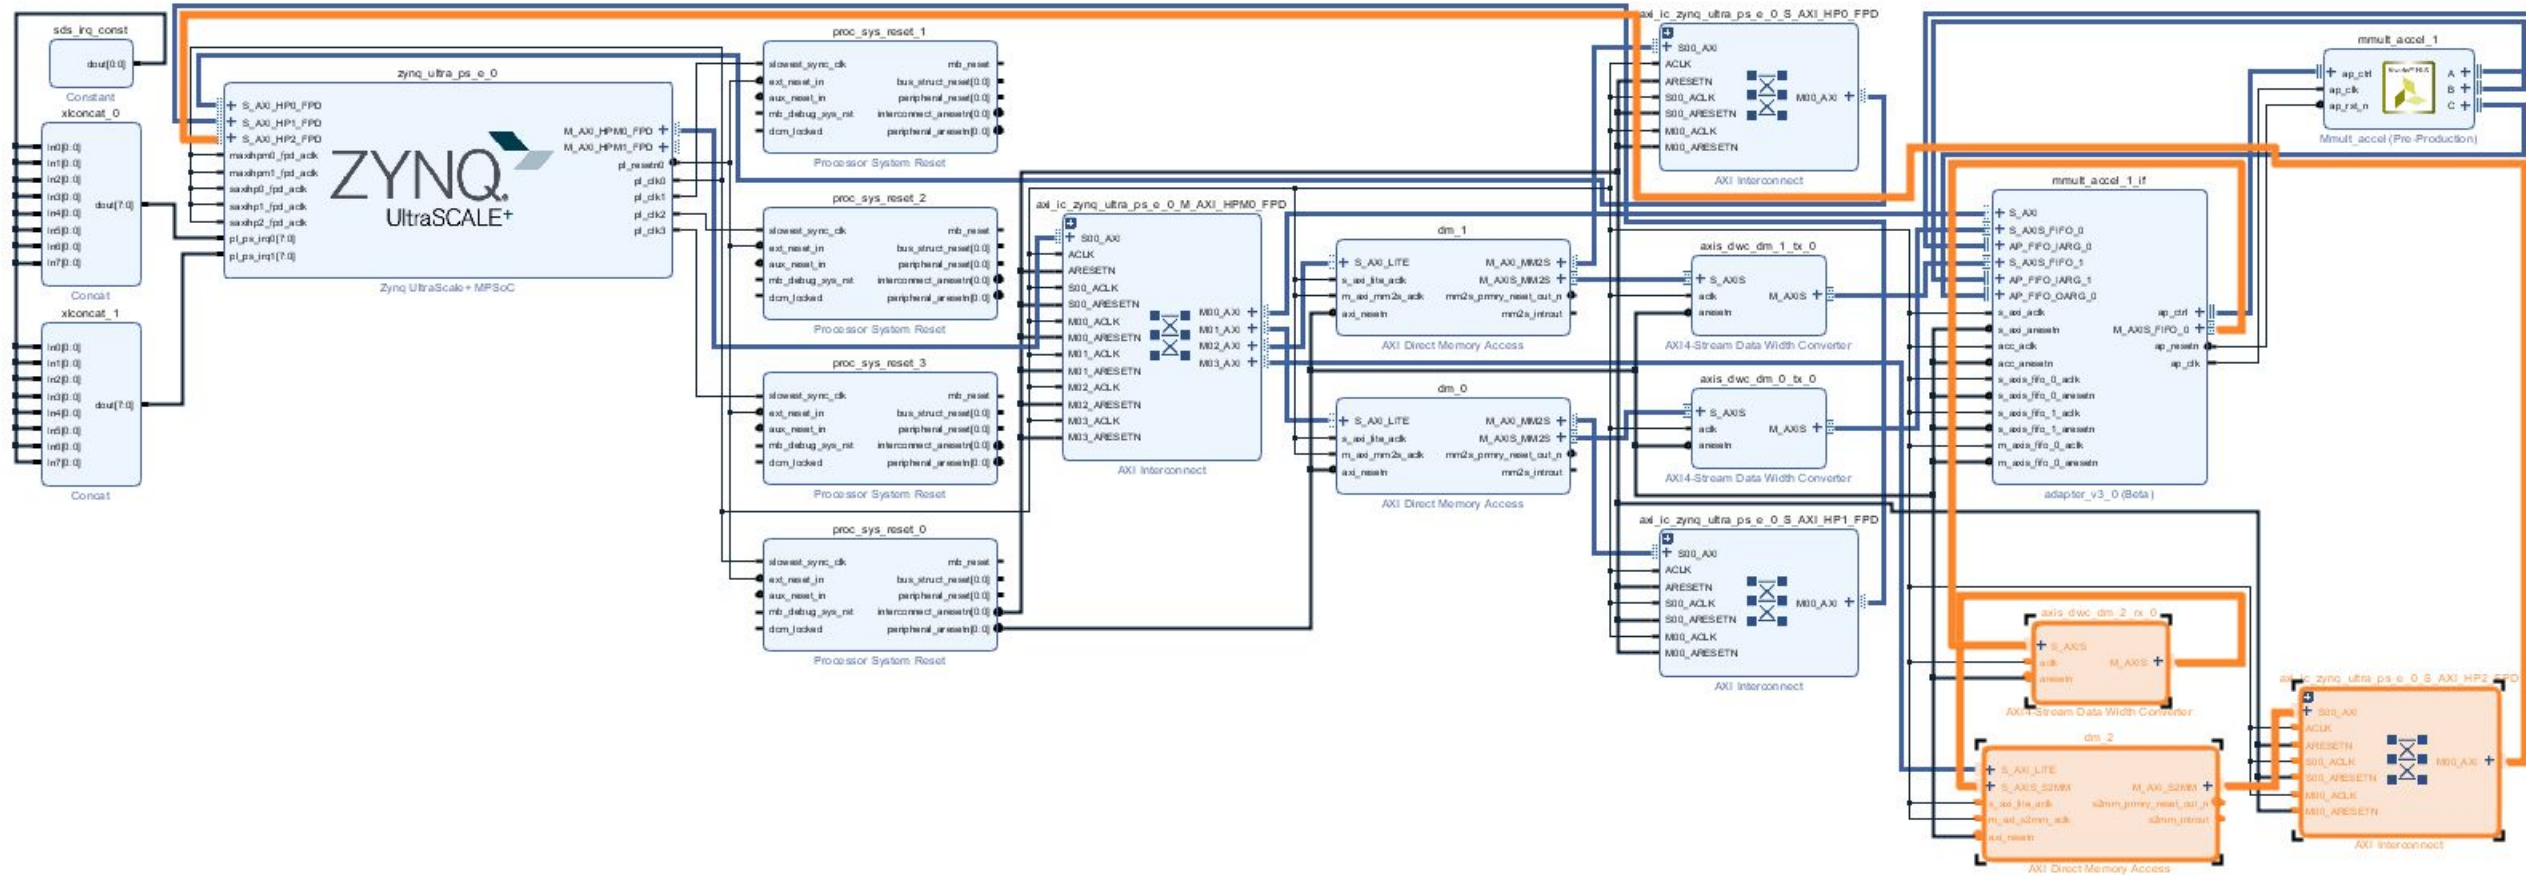
## Example: Ultra96 Platform with mmult



# Example: Ultra96 Platform with mmult



## Example: Ultra96 Platform with mmult



Thank you!

# Features



---

# C-Callable IP

---

## C-Callable IP

- Allows the reuse legacy IP and provide a software-centric handoff from hardware / IP team to application developers
- Launch from SDx Terminal

# Additional Resources

# For more Information on HLS or Pragmas

---

## SDx Developer Topics

- [https://www.xilinx.com/html\\_docs/xilinx2017\\_2/sdaccel\\_doc/topics/pragmas/ref-pragma\\_HLS\\_interface.html](https://www.xilinx.com/html_docs/xilinx2017_2/sdaccel_doc/topics/pragmas/ref-pragma_HLS_interface.html)

## Getting Started with Vivado High-Level Synthesis

- <https://www.xilinx.com/video/hardware/getting-started-vivado-high-level-synthesis.html>
- Avnet's SDSoC Advanced Concepts – Streaming IO Example
  - Please visit the minized.org website or contact your local Avnet FAE for access
  - This follows the UG1146 Example injecting streaming IO control into an SDSoC Platform
  - Based on our MiniZed Platform

---

# For more Information on HLS or Pragmas

---

Xilinx Documentation (not all inclusive)

- Vivado Design Suite Tutorial High-Level Synthesis UG871 (v2017.4)
- Vivado Design Suite User Guide High-Level Synthesis UG902 (v2017.4)
- SDSoC Environment User Guide UG1027 (v2017.4)
- SDx Pragma Reference Guide UG1253 (v2017.4)

# Basic HLS

# HLS

---

## High Level Synthesis

- Vivado High-Level Synthesis compiler
  - C, C++ and SystemC programs directly targeted into Xilinx devices
  - Without the need to manually create RTL

[https://en.wikipedia.org/wiki/Xilinx\\_Vivado](https://en.wikipedia.org/wiki/Xilinx_Vivado)

- What does this mean?
  - You can write HLS compliant C/C++ code and the TOOL will translate that to RTL!

---

# How does it know what to do?

---

## SDSoC Uses

- Modeling
- Estimation
- Predictive Analysis

Through this, SDSoC can

- Utilize all available resources to create the most optimal design given the resources available

# How can I tell the tool what I want?

---

- SDSoC Leverages a Suite of Tools, one being Vivado HLS pragmas!
- In C based design, all input and output operations are performed, in zero time, through formal function arguments.
- In an RTL design these same input and output operations must be performed through a port in the design interface and typically operate using a specific I/O (input-output) protocol.
- For more information, refer to "Managing Interfaces" in the Vivado Design Suite User Guide: High-Level Synthesis (UG902).



---

# How can I tell the tool what I want?

---

In other words

- Pragmas are comments placed into the C/C++ code that instructs the C->RTL conversion to have specific bounds and goals
- Remember, SDSoC sits on TOP of a suite of tools
  - While there are HLS specific pragmas, there are also SDS pragmas to be used with the SDx family

# Pragmas Use In Our Example

---

mmult.h

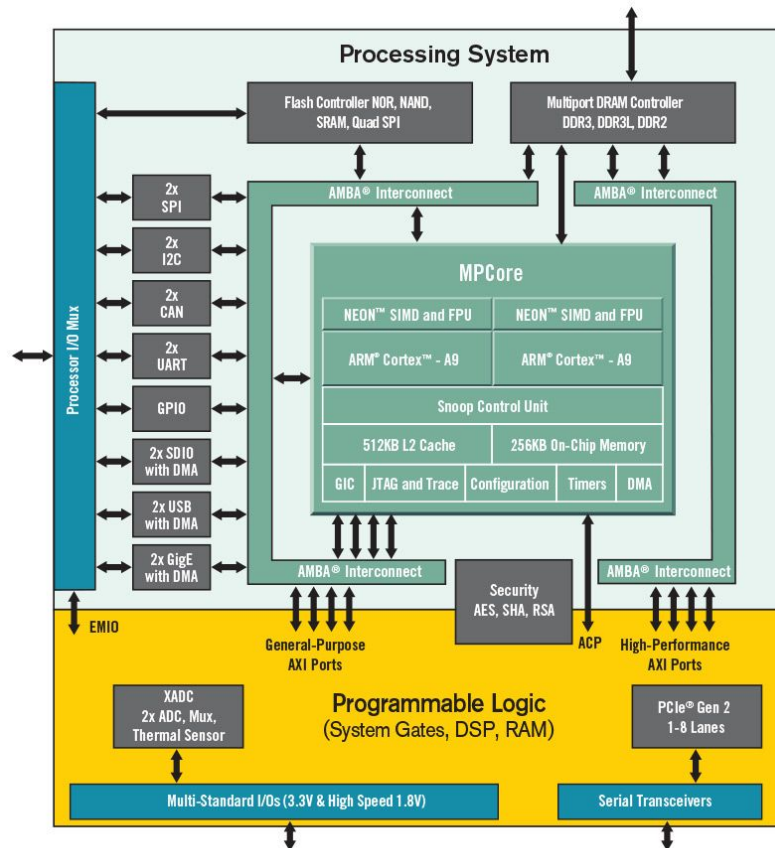
- `#pragma SDS data access_pattern`
  - This pragma specifies the data access pattern in the hardware function.
  - By describing the KNOWN data access pattern, SDSoC can best determine what data movers should be used

mmult.cpp

- `#pragma HLS array_partition`
  - Partitions an array into smaller arrays or individual elements.
  - In this case used to spread out array allowing parallel access to all elements
- `#pragma HLS PIPELINE`
  - The PIPELINE pragma reduces the initiation interval for a function or loop by allowing the concurrent execution of operations.
  - Adding this will generate a pipelined feeder to ensure new data is available to the function or look every N clock cycles
  - In this case every 1 clock cycle

\*information from Xilinx SDx Developer Topics website

# What is Zynq?



- Single or Dual Core Arm A9
- NEON SIMD Accelerators
- Variable amounts of Programmable Logic (amount of acceleration possible)
- Variable grades of Programmable Logic (speed)
- Variable fabric families of Programmable Logic (performance)

# What is Zynq?



- EX: Avnet MiniZed \$88.99; comes with SDSoc License
- Arduino Compatible